

Lab Bayesian

Andres Navarro

2023-04-25

Dataset Lookup

First of all, the objective of this Case Study will be to show the advantages of Bayesian Statistics for small data sets and the ability to estimated better the posterior parameters.. As it is known, Bayesian Statistics allows as to set up prior believes of our predictors with specific probability distributions. This is really useful when we do not have a lot of data and we have some insights on the data. For this reason I have decided to use a data set with 21 variables and I will be reducing the number of observations to simulate what we are trying to show. This data is about the COVID cases in Mexico and the goal is to predict if a patient has COVID or not.

[[<https://www.kaggle.com/datasets/meirnazri/covid19-dataset>][Dataset]]

```
rm(list = ls())
data = read.csv("data.csv", header = TRUE)
dim(data)
```

```
## [1] 1048575      21
```

```
summary(data)
```

```
##      USMER      MEDICAL_UNIT      SEX      PATIENT_TYPE
## Min.   :1.000   Min.   : 1.000   Min.   :1.000   Min.   :1.000
## 1st Qu.:1.000   1st Qu.: 4.000   1st Qu.:1.000   1st Qu.:1.000
## Median :2.000   Median :12.000   Median :1.000   Median :1.000
## Mean   :1.632   Mean   : 8.981   Mean   :1.499   Mean   :1.191
## 3rd Qu.:2.000   3rd Qu.:12.000   3rd Qu.:2.000   3rd Qu.:1.000
## Max.   :2.000   Max.   :13.000   Max.   :2.000   Max.   :2.000
##      DATE_DIED      INTUBED      PNEUMONIA      AGE
## Length:1048575   Min.   : 1.00   Min.   : 1.000   Min.   : 0.00
## Class :character 1st Qu.:97.00 1st Qu.: 2.000   1st Qu.: 30.00
## Mode  :character Median :97.00 Median : 2.000   Median : 40.00
##                Mean   :79.52 Mean   : 3.347   Mean   : 41.79
##                3rd Qu.:97.00 3rd Qu.: 2.000   3rd Qu.: 53.00
##                Max.   :99.00 Max.   :99.000   Max.   :121.00
##      PREGNANT      DIABETES      COPD      ASTHMA
## Min.   : 1.00   Min.   : 1.000   Min.   : 1.000   Min.   : 1.000
## 1st Qu.: 2.00   1st Qu.: 2.000   1st Qu.: 2.000   1st Qu.: 2.000
## Median :97.00   Median : 2.000   Median : 2.000   Median : 2.000
## Mean   :49.77   Mean   : 2.186   Mean   : 2.261   Mean   : 2.243
## 3rd Qu.:97.00   3rd Qu.: 2.000   3rd Qu.: 2.000   3rd Qu.: 2.000
## Max.   :98.00   Max.   :98.000   Max.   :98.000   Max.   :98.000
##      INMSUPR      HIPERTENSION      OTHER_DISEASE      CARDIOVASCULAR
## Min.   : 1.000   Min.   : 1.000   Min.   : 1.000   Min.   : 1.000
## 1st Qu.: 2.000   1st Qu.: 2.000   1st Qu.: 2.000   1st Qu.: 2.000
```

```
## Median : 2.000 Median : 2.000 Median : 2.000 Median : 2.000
## Mean : 2.298 Mean : 2.129 Mean : 2.435 Mean : 2.262
## 3rd Qu.: 2.000 3rd Qu.: 2.000 3rd Qu.: 2.000 3rd Qu.: 2.000
## Max. :98.000 Max. :98.000 Max. :98.000 Max. :98.000
## OBESITY RENAL_CHRONIC TOBACCO CLASIFFICATION_FINAL
## Min. : 1.000 Min. : 1.000 Min. : 1.000 Min. :1.000
## 1st Qu.: 2.000 1st Qu.: 2.000 1st Qu.: 2.000 1st Qu.:3.000
## Median : 2.000 Median : 2.000 Median : 2.000 Median :6.000
## Mean : 2.125 Mean : 2.257 Mean : 2.214 Mean :5.306
## 3rd Qu.: 2.000 3rd Qu.: 2.000 3rd Qu.: 2.000 3rd Qu.:7.000
## Max. :98.000 Max. :98.000 Max. :98.000 Max. :7.000
## ICU
## Min. : 1.00
## 1st Qu.:97.00
## Median :97.00
## Mean :79.55
## 3rd Qu.:97.00
## Max. :99.00
```

The raw data set consists of 21 unique features and 1,048,576 unique patients. In the Boolean features, 1 means “yes” and 2 means “no”. values as 97 and 99 are missing data.

- sex: 1 for female and 2 for male.
- age: of the patient.
- classification: covid test findings. Values 1-3 mean that the patient was diagnosed with covid in different
- degrees. 4 or higher means that the patient is not a carrier of covid or that the test is inconclusive.
- patient type: type of care the patient received in the unit. 1 for returned home and 2 for hospitalization.
- pneumonia: whether the patient already have air sacs inflammation or not.
- pregnancy: whether the patient is pregnant or not.
- diabetes: whether the patient has diabetes or not.
- copd: Indicates whether the patient has Chronic obstructive pulmonary disease or not.
- asthma: whether the patient has asthma or not.
- inmsupr: whether the patient is immunosuppressed or not.
- hypertension: whether the patient has hypertension or not.
- cardiovascular: whether the patient has heart or blood vessels related disease.
- renal chronic: whether the patient has chronic renal disease or not.
- other disease: whether the patient has other disease or not.
- obesity: whether the patient is obese or not.
- tobacco: whether the patient is a tobacco user.
- usmr: Indicates whether the patient treated medical units of the first, second or third level.
- medical unit: type of institution of the National Health System that provided the care.
- intubed: whether the patient was connected to the ventilator.
- icu: Indicates whether the patient had been admitted to an Intensive Care Unit.
- date died: If the patient died indicate the date of death, and 9999-99-99 otherwise.

Here we can see a summary of the data, first we have to clean and adapt the data so we can work on it. First of all, I will create the variable that we want to predict that is if a patient has been diagnosed with COVID or not.

```
data$COVID = ifelse(data$CLASIFFICATION_FINAL <= 3, 1, 2)
data = subset(data, select = -c(CLASIFFICATION_FINAL))
```

```
convertToLogic = function(col.name, df) {
  index = which(names(df) == col.name)
  print(index)
```

```

if (length(index) != 0) {
  df[, index] = ifelse(df[, index] == 2, 0, df[, index])
  df[, index] = as.logical(df[, index])
}

return(df)
}

```

This column will tell us if a patient has been diagnosed with COVID or not. Then, I will factor and format all the other variables to adapt them properly.

```
data = convertToLogic("COVID", data)
```

```
## [1] 21
```

```

data$USMER = ifelse(data$USMER == 2, 0, data$USMER)
data$USMER = as.logical(data$USMER)

data$MEDICAL_UNIT = factor(data$MEDICAL_UNIT)

data$SEX = factor(data$SEX, labels = c("female", "male"), levels = c(1, 2))

data$PATIENT_TYPE = factor(data$PATIENT_TYPE, labels = c("returned home", "hospitalized"), levels = c(1, 2))

data$INTUBED = factor(data$INTUBED, labels = c("intubed", "not intubed"), levels = c(1, 2))

data$PNEUMONIA = factor(data$PNEUMONIA, labels = c("pneumonia", "not pneumonia"), levels = c(1, 2))

data$PREGNANT = factor(data$PREGNANT, labels = c("pregnant", "not pregnant"), levels = c(1, 2))

data$DIABETES = factor(data$DIABETES, labels = c("diabetes", "not diabetes"), levels = c(1, 2))

data$COPD = factor(data$COPD, labels = c("copd", "not copd"), levels = c(1, 2))

data$ASTHMA = factor(data$ASTHMA, labels = c("asthma", "not asthma"), levels = c(1, 2))

data$INMSUPR = factor(data$INMSUPR, labels = c("inmsupr", "not inmsupr"), levels = c(1, 2))

data$HIPERTENSION = factor(data$HIPERTENSION, labels = c("hipertension", "not hipertension"), levels = c(1, 2))

data$OTHER_DISEASE = factor(data$OTHER_DISEASE, labels = c("other disease", "not other disease"), levels = c(1, 2))

data$CARDIOVASCULAR = factor(data$CARDIOVASCULAR, labels = c("cardiovascular", "not cardiovascular"), levels = c(1, 2))

data$OBESITY = factor(data$OBESITY, labels = c("obesity", "not obesity"), levels = c(1, 2))

data$RENAL_CHRONIC = factor(data$RENAL_CHRONIC, labels = c("renal chronic", "not renal chronic"), levels = c(1, 2))

data$TOBACCO = factor(data$TOBACCO, labels = c("tobacco", "not tobacco"), levels = c(1, 2))

data$ICU = factor(data$ICU, labels = c("icu", "not icu"), levels = c(1, 2))

data = subset(data, select = -c(DATE_DIED))

```

```
summary(data)
```

```
##      USMER      MEDICAL_UNIT      SEX      PATIENT_TYPE
## Mode :logical  12      :602995  female:525064  returned home:848544
## FALSE:662903  4      :314405  male :523511  hospitalized :200031
## TRUE :385672  6      : 40584
##      9      : 38116
##      3      : 19175
##      8      : 10399
##      (Other): 22901
##      INTUBED      PNEUMONIA      AGE
## intubed   : 33656  pneumonia :140038  Min.   : 0.00
## not intubed:159050 not pneumonia:892534 1st Qu.: 30.00
## NA's      :855869  NA's      : 16003  Median : 40.00
##                                     Mean    : 41.79
##                                     3rd Qu.: 53.00
##                                     Max.    :121.00
##
##      PREGNANT      DIABETES      COPD
## pregnant   : 8131  diabetes   :124989  copd    : 15062
## not pregnant:513179 not diabetes:920248 not copd:1030510
## NA's       :527265  NA's       : 3338  NA's    : 3003
##
##
##
##      ASTHMA      INMSUPR      HIPERTENSION
## asthma      : 31572  inmsupr    : 14170  hypertension :162729
## not asthma:1014024 not inmsupr:1031001 not hypertension:882742
## NA's        : 2979  NA's        : 3404  NA's        : 3104
##
##
##
##      OTHER_DISEASE      CARDIOVASCULAR      OBESITY
## other disease : 28040  cardiovascular : 20769  obesity :159816
## not other disease:1015490 not cardiovascular:1024730 not obesity:885727
## NA's          : 5045  NA's          : 3076  NA's      : 3032
##
##
##
##      RENAL_CHRONIC      TOBACCO      ICU
## renal chronic : 18904  tobacco : 84376  icu : 16858
## not renal chronic:1026665 not tobacco:960979 not icu:175685
## NA's          : 3006  NA's      : 3220  NA's      :856032
##
##
##
##      COVID
## Mode :logical
## FALSE:656596
## TRUE :391979
```

```
##  
##  
##  
##
```

Here we see that the data is correctly formatted but there are some missing value, so let's fix that.

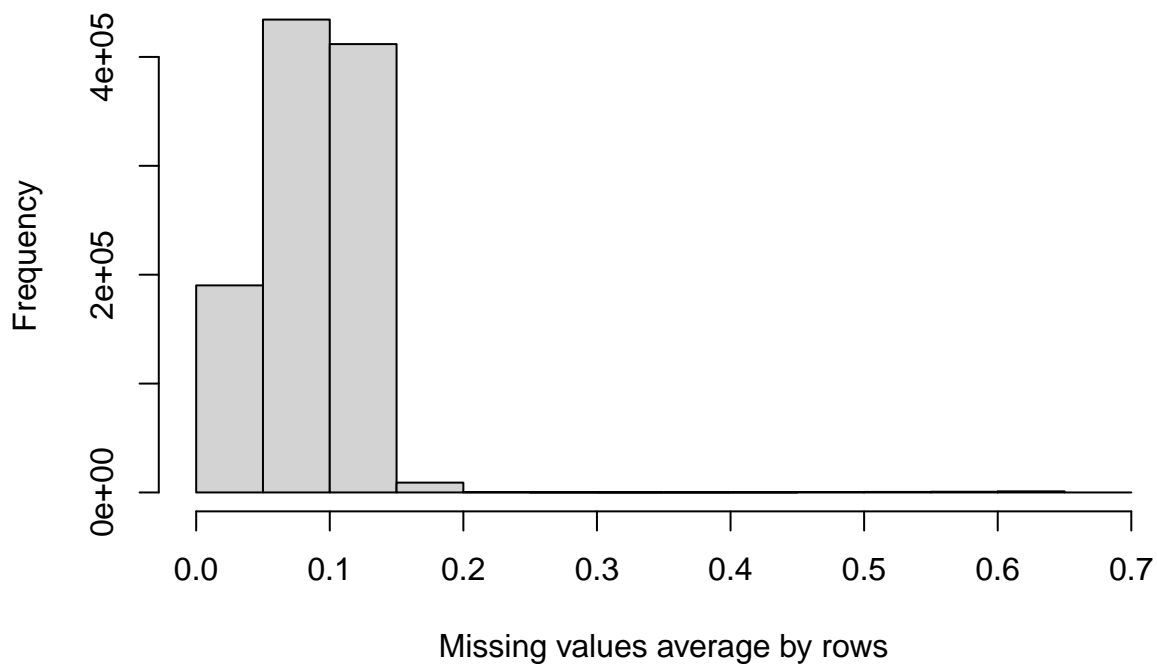
Data Cleaning

First we will see how many missing values there are by rows so we can remove some columns that have a lot of missing values.

```
print(length(which(is.na(data))))
```

```
## [1] 2288376
```

```
hist(rowMeans(is.na(data)), xlab = c("Missing values average by rows"), main = c())
```




Here we see that there are 3 columns with the most missing values.

```
indexesEmptyCols = which(colMeans(is.na(data)) != 0)
```

```
colsWithNA = sort(colMeans(is.na(data[, indexesEmptyCols])),  
                  decreasing = TRUE)
```

```
barplot(colsWithNA, las=2)
```



And the columns that have the most missing values are ICU, INTUBED, and PREGNANT, so let's remove them.

```
data = subset(data, select = -c(ICU, INTUBED, PREGNANT))
```

```
print(length(which(is.na(data))))
```

```
## [1] 49210
```

```
data = na.omit(data)
```

```
length(unique(which(is.na(data))))
```

```
## [1] 0
```

```
summary(data)
```

```
##      USMER      MEDICAL_UNIT      SEX      PATIENT_TYPE
## Mode :logical  12      :591811  female:513216  returned home:833253
## FALSE:658255   4      :307177  male  :511936  hospitalized :191899
## TRUE :366897   6      : 37868
##                      9      : 37384
##                      3      : 18660
##                      8      : 10097
##                      (Other): 22155
##      PNEUMONIA      AGE      DIABETES
## pneumonia      :137599  Min.   : 0.00  diabetes      :122415
## not pneumonia:887553  1st Qu.: 30.00  not diabetes:902737
##                      Median : 40.00
##                      Mean   : 41.89
##                      3rd Qu.: 53.00
##                      Max.   :121.00
##
```

```
##          COPD          ASTHMA          INMSUPR
##  copd      : 14376   asthma      : 30497   inmsupr      : 13588
##  not copd:1010776   not asthma:994655   not inmsupr:1011564
##
##
##
##
##          HIPERTENSION          OTHER_DISEASE
##  hypertension      :159577   other disease      : 27131
##  not hypertension:865575   not other disease:998021
##
##
##
##
##          CARDIOVASCULAR          OBESITY          RENAL_CHRONIC
##  cardiovascular      : 20126   obesity      :156961   renal chronic      : 18351
##  not cardiovascular:1005026   not obesity:868191   not renal chronic:1006801
##
##
##
##
##          TOBACCO          COVID
##  tobacco      : 82675   Mode :logical
##  not tobacco:942477   FALSE:636274
##                      TRUE :388878
##
##
##
##
```

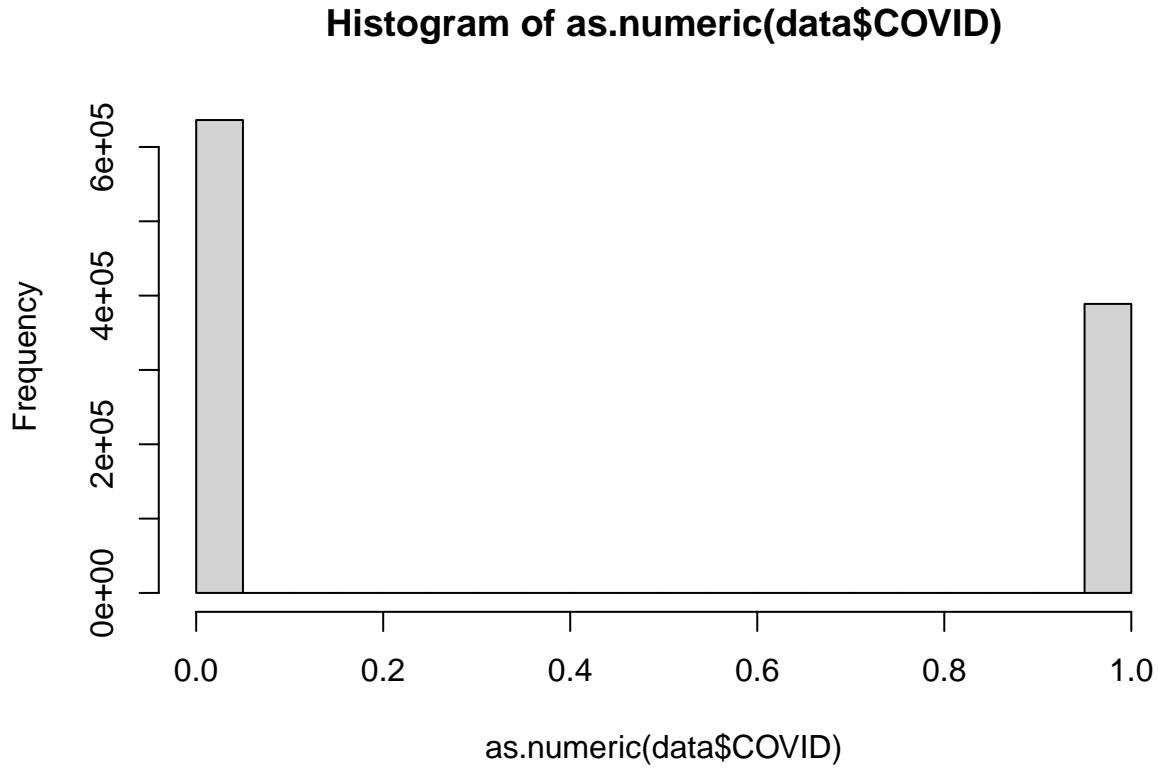
Now the data set is clean so let's start working on it. But first we will shrink it to 1000 observations to work with.

```
data.small = data[sample(nrow(data), size=1000),]
```

Bayesian Analysis of the covid variable

First of all, let's plot a histogram of the COVID variable (the one we want to predict) and see.

```
rm(list = setdiff(ls(), c("data", "data.small")))
hist(as.numeric(data$COVID))
```



This is as we expected as we are going to be predicting a binary variable. So let's use a Bernoulli distribution to explain this data and see how well it fits. First of all let's compute the analytical posterior distribution of the covid variable.

Analytical Study

1. We assume a Bernoulli distribution for COVID, we will use X to denote that variable.

$$X \mid \theta \sim \text{Bernoulli}(\theta)$$

$$f(x \mid \theta) = \theta^x \cdot (1 - \theta)^{1-x}$$

2. As we do not have any prior knowledge on the probability of a patient of having covid, we will define the prior distribution as an improper prior. Moreover, we will be using a Beta distribution as in the end we will get a posterior conjugate which will be much easier to work with.

$$\theta \sim \text{Beta}(0, 0)$$

$$f(\theta \mid 0, 0) = \frac{\theta^{0-1} \cdot (1 - \theta)^{0-1}}{B(0, 0)}$$

3. Now we get the likelihood

$$f(\text{data} \mid \theta) \propto \theta^k \cdot (1 - \theta)^{n-k}$$

Being n the total number of observations and k the positive ones.

4. And finally the posterior distribution

$$f(\theta | data) = \frac{\theta^{k-1} \cdot (1-\theta)^{n-k-1}}{B(k, n-k)}$$

$$\theta | data \sim Beta(k, n-k)$$

So now that we have the posterior distribution let's obtain the prediction of the next value called Y given the data

$$Y | \theta \sim Bernuilli(\theta)$$

$$P(Y = 1|data) = \int_{-\infty}^{\infty} P(Y = 1|\theta) \cdot P(\theta|data)d\theta = \frac{B(k+1, n-k)}{B(k, n-k)}$$

```
n = as.numeric(length(data.small$COVID))
k = as.numeric(length(which(data.small$COVID)))
print(beta(k+1,n-k)/beta(k, n-k))
```

```
## [1] 0.387
```

And here we can see that the probability of a new patient of having covid is 0.362 that is really close to the ML estimator of 0.38

And finally let's try to obtain the same result numerically

Numerical Study

As we know the distribution of the new observation we will get a random sample and compare.

$$Y | \theta \sim Bernuilli(Beta(k, n-k))$$

```
y.sample = rbinom(n, 1, rbeta(1, k, n - k))
mean(y.sample)
```

```
## [1] 0.397
```

Here we see that the estimated probability is almost the same as previously.

```
covid.prob = rbeta(n, k, n - k)
quantile(covid.prob, probs = c(0.025, 0.975))
```

```
##      2.5%      97.5%
## 0.3566339 0.4157329
```

And also we see that the confidence interval for the probability of having covid is pretty narrow, so we can be sure that it is correct.


Data Exploration

Now, we will see if the other variables are useful to predict if a patient has covid or not.

```
rm(list = setdiff(ls(), c("data", "data.small")))

library(ggplot2) # GGally
library(GGally)
```

```
## Registered S3 method overwritten by 'GGally':
##   method from
##   +.gg      ggplot2
ggcorr(data, cor_matrix = cor(sapply(data, as.numeric)), label = TRUE)
```




Here we see that there is some correlation between the columns but nothing strong with respect to the COVID variable so let's see if we can identify better which columns have more correlation with the covid column.

```
corr_covid = sort(cor(sapply(subset(data, select = -c(COVID)), as.numeric))[1,], decreasing = T)

corr = data.frame(corr_covid)


ggplot(corr, aes(x = row.names(corr), y = corr_covid)) + geom_bar(stat = "identity") +
  scale_x_discrete(limits = row.names(corr)) + labs(x = "", y = "covid", title = "Correlations") +
  theme(plot.title = element_text(hjust = 0, size = rel(1.5)), axis.text.x = element_text(angle = 45, hjust = 1))
```



Here we see that USMER has the most correlation and this makes sense as it indicated if the patient has received medication.

Now let's see how the columns distribute with respect to the covid variable.


```
ggplot(data, aes(x=COVID, y=AGE)) +  
  geom_boxplot()
```



Here we see that there is a visible difference between the mean of the covid, so this can be a useful variable to use in our model.


```
plot(table(data$COVID, data$USMER))
```

table(data\$COVID, data\$USMER)




```
plot(table(data$COVID, data$MEDICAL_UNIT))
```

table(data\$COVID, data\$MEDICAL_UNIT)




```
plot(table(data$COVID, data$SEX))
```

table(data\$COVID, data\$SEX)




```
plot(table(data$COVID, data$PATIENT_TYPE))
```

table(data\$COVID, data\$PATIENT_TYPE)




```
plot(table(data$COVID, data$PNEUMONIA))
```

table(data\$COVID, data\$PNEUMONIA)




```
plot(table(data$COVID, data$DIABETES))
```

table(data\$COVID, data\$DIABETES)




```
plot(table(data$COVID, data$COPD))
```

table(data\$COVID, data\$COPD)




```
plot(table(data$COVID, data$ASTHMA))
```

table(data\$COVID, data\$ASTHMA)




```
plot(table(data$COVID, data$INMSUPR))
```

table(data\$COVID, data\$INMSUPR)




```
plot(table(data$COVID, data$HIPERTENSION))
```


table(data\$COVID, data\$HIPERTENSION)




```
plot(table(data$COVID, data$OTHER_DISEASE))
```

table(data\$COVID, data\$OTHER_DISEASE)




```
plot(table(data$COVID, data$CARDIOVASCULAR))
```

table(data\$COVID, data\$CARDIOVASCULAR)




```
plot(table(data$COVID, data$OBESITY))
```

table(data\$COVID, data\$OBESITY)




```
plot(table(data$COVID, data$RENAL_CHRONIC))
```

```
table(data$COVID, data$RENAL_CHRONIC)
```



```
plot(table(data$COVID, data$TOBACCO))
```

```
table(data$COVID, data$TOBACCO)
```



Now, the columns that show off the most are: MEDICAL_UNIT, SEX, PATIENT_TYPE, and PNEUMONIA. This makes sense and we will see after if we are confident that there is a visible difference.

Frequentist LM

Now let's implement a simple LM model to see how well we can predict a patient to have covid.

```
rm(list = setdiff(ls(), c("data")))
```

```
library(caret)
```

```
## Loading required package: lattice
```

```
library(lattice)

data.small = data[sample(nrow(data), size=10000),]

index.test = createDataPartition(data.small$COVID, p = 0.5, list = FALSE)

data.test = data.small[index.test,]
data.train = data.small[-index.test,]

rm(index.test)
```

Now first of all let's try to use all the variables to try to predict if a patient has covid or not.

```
fit = train(as.factor(COVID) ~ ., data = data.train, method = "glm", family = "binomial")

summary(fit)
```

```
##
## Call:
## NULL
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -1.549007    1.040384  -1.489  0.136519
## USMERTRUE      -0.006100    0.065788  -0.093  0.926121
## MEDICAL_UNIT2 -11.032253  196.969748  -0.056  0.955334
## MEDICAL_UNIT3    0.661634    0.925289   0.715  0.474574
## MEDICAL_UNIT4    0.372106    0.896533   0.415  0.678106
## MEDICAL_UNIT5    0.579578    0.968971   0.598  0.549748
## MEDICAL_UNIT6   -0.083100    0.909352  -0.091  0.927188
## MEDICAL_UNIT7   -0.885892    1.491755  -0.594  0.552606
## MEDICAL_UNIT8    0.403198    0.954278   0.423  0.672648
## MEDICAL_UNIT9    0.563616    0.909098   0.620  0.535276
## MEDICAL_UNIT10   1.173119    0.962655   1.219  0.222985
## MEDICAL_UNIT11   0.685636    0.975880   0.703  0.482316
## MEDICAL_UNIT12   0.295849    0.896320   0.330  0.741347
## MEDICAL_UNIT13           NA           NA           NA           NA
## SEXmale          0.225557    0.061613   3.661  0.000251 ***
## PATIENT_TYPEhospitalized 0.517222    0.105990   4.880  1.06e-06 ***
## `PNEUMONIAnot pneumonia` -0.487553    0.116301  -4.192  2.76e-05 ***
## AGE              0.009648    0.002100   4.593  4.36e-06 ***
## `DIABETESnot diabetes`   -0.283205    0.104088  -2.721  0.006512 **
## `COPDnot copd`          0.230603    0.276268   0.835  0.403882
## `ASTHMAnot asthma`       0.151850    0.185813   0.817  0.413805
## `INMSUPRnot inmsupr`     0.016963    0.265770   0.064  0.949110
## `HIPERTENSIONnot hypertension` -0.103733    0.096202  -1.078  0.280911
## `OTHER_DISEASEnot other disease` 0.234106    0.179897   1.301  0.193144
## `CARDIOVASCULARnot cardiovascular` 0.268790    0.225357   1.193  0.232977
## `OBESITYnot obesity`     -0.425641    0.086943  -4.896  9.80e-07 ***
## `RENAL_CHRONICnot renal chronic`  0.243150    0.245701   0.990  0.322361
## `TOBACCOnot tobacco`      0.056725    0.112339   0.505  0.613596
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

```
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 6562.6 on 4998 degrees of freedom
## Residual deviance: 6247.4 on 4972 degrees of freedom
## AIC: 6301.4
##
## Number of Fisher Scoring iterations: 10
```

Here we see that there are a lot of variables that are useless. As the p value of the betas is really high for most of them.

```
confusionMatrix(as.factor(data.test$COVID), predict(fit, newdata = data.test))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction FALSE TRUE
##      FALSE  2851  323
##      TRUE   1340  487
##
##           Accuracy : 0.6675
##           95% CI : (0.6542, 0.6805)
##      No Information Rate : 0.838
##      P-Value [Acc > NIR] : 1
##
##           Kappa : 0.1869
##
## Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.6803
##           Specificity : 0.6012
##           Pos Pred Value : 0.8982
##           Neg Pred Value : 0.2666
##           Prevalence : 0.8380
##           Detection Rate : 0.5701
##           Detection Prevalence : 0.6347
##           Balanced Accuracy : 0.6408
##
##           'Positive' Class : FALSE
##
```

Here, we see that we get an accuracy of 0.6578 so it is not that bad, probably it is because we only have a few significant variables as we saw in the correlation graph. So let's try a simpler model.

```
fit = train(as.factor(COVID) ~ USMER + PNEUMONIA + MEDICAL_UNIT + DIABETES + HIPERTENSION + AGE + PATIENT)
summary(fit)
```

```
##
## Call:
## NULL
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -0.693216   0.930881  -0.745   0.4565
## USMERTRUE    -0.014868   0.065410  -0.227   0.8202
```

```

## `PNEUMONIA`not pneumonia`      -0.521051   0.115291  -4.519 6.20e-06 ***
## MEDICAL_UNIT2                    -10.807228 196.969791 -0.055 0.9562
## MEDICAL_UNIT3                     0.807632   0.934440   0.864 0.3874
## MEDICAL_UNIT4                     0.491514   0.906614   0.542 0.5877
## MEDICAL_UNIT5                     0.741058   0.977544   0.758 0.4484
## MEDICAL_UNIT6                     0.030421   0.919609   0.033 0.9736
## MEDICAL_UNIT7                    -0.568439   1.468192  -0.387 0.6986
## MEDICAL_UNIT8                     0.721489   0.961007   0.751 0.4528
## MEDICAL_UNIT9                     0.704366   0.918767   0.767 0.4433
## MEDICAL_UNIT10                    1.296947   0.972232   1.334 0.1822
## MEDICAL_UNIT11                    0.855189   0.986170   0.867 0.3858
## MEDICAL_UNIT12                    0.426837   0.906500   0.471 0.6377
## MEDICAL_UNIT13                     NA         NA         NA         NA
## `DIABETES`not diabetes`          -0.260212   0.102917  -2.528 0.0115 *
## `HIPERTENSION`not hipertension`  -0.140968   0.094053  -1.499 0.1339
## AGE                               0.009311   0.002071   4.496 6.91e-06 ***
## PATIENT_TYPEhospitalized         0.518198   0.104440   4.962 6.99e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 6562.6  on 4998  degrees of freedom
## Residual deviance: 6290.2  on 4981  degrees of freedom
## AIC: 6326.2
##
## Number of Fisher Scoring iterations: 10

```

Now it is better but the medical unit for example, it is only relevant the level 2 and also for other.

```
confusionMatrix(as.factor(data.test$COVID), predict(fit, newdata = data.test))
```

```

## Confusion Matrix and Statistics
##
##              Reference
## Prediction FALSE TRUE
##      FALSE  2857  317
##      TRUE   1356  471
##
##              Accuracy : 0.6655
##              95% CI : (0.6522, 0.6785)
##      No Information Rate : 0.8424
##      P-Value [Acc > NIR] : 1
##
##              Kappa : 0.1796
##
##      Mcnemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.6781
##              Specificity : 0.5977
##      Pos Pred Value : 0.9001
##      Neg Pred Value : 0.2578
##              Prevalence : 0.8424
##      Detection Rate : 0.5713
##      Detection Prevalence : 0.6347

```

```
##      Balanced Accuracy : 0.6379
##
##      'Positive' Class : FALSE
##
```

Here we see that the accuracy is almost the same and the kappa so we have not lost a lot of info.

Bayesian LM

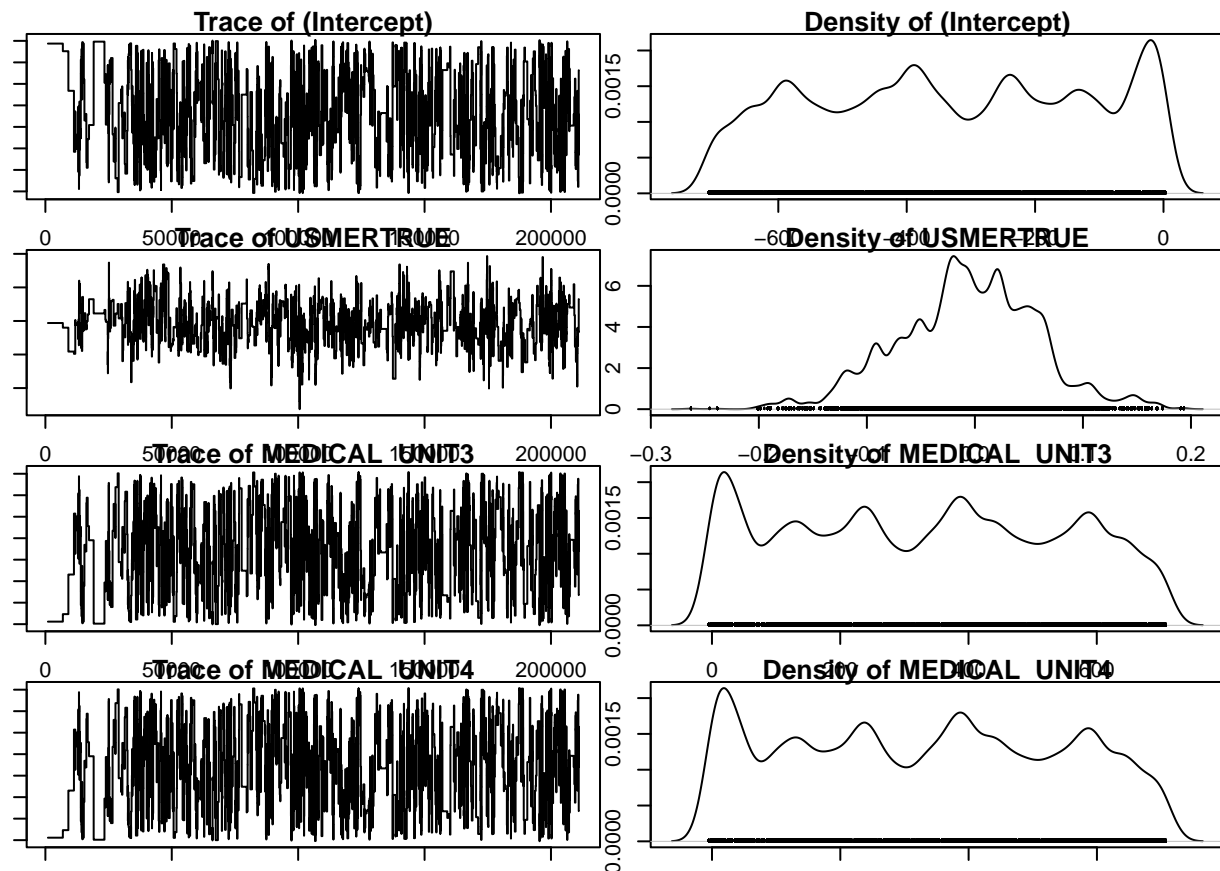
The frequentist approach is easier but we if we want to compute confidence intervals for the parameters or predictive intervals we cannot do them. That is why we will be using the Bayesian approach to better study the effects of each variable with covid and get more conclusions. The power of the Bayesian approach is that we obtain the posterior distribution of the parameters so we can study better the relation and the significance. So let's start.

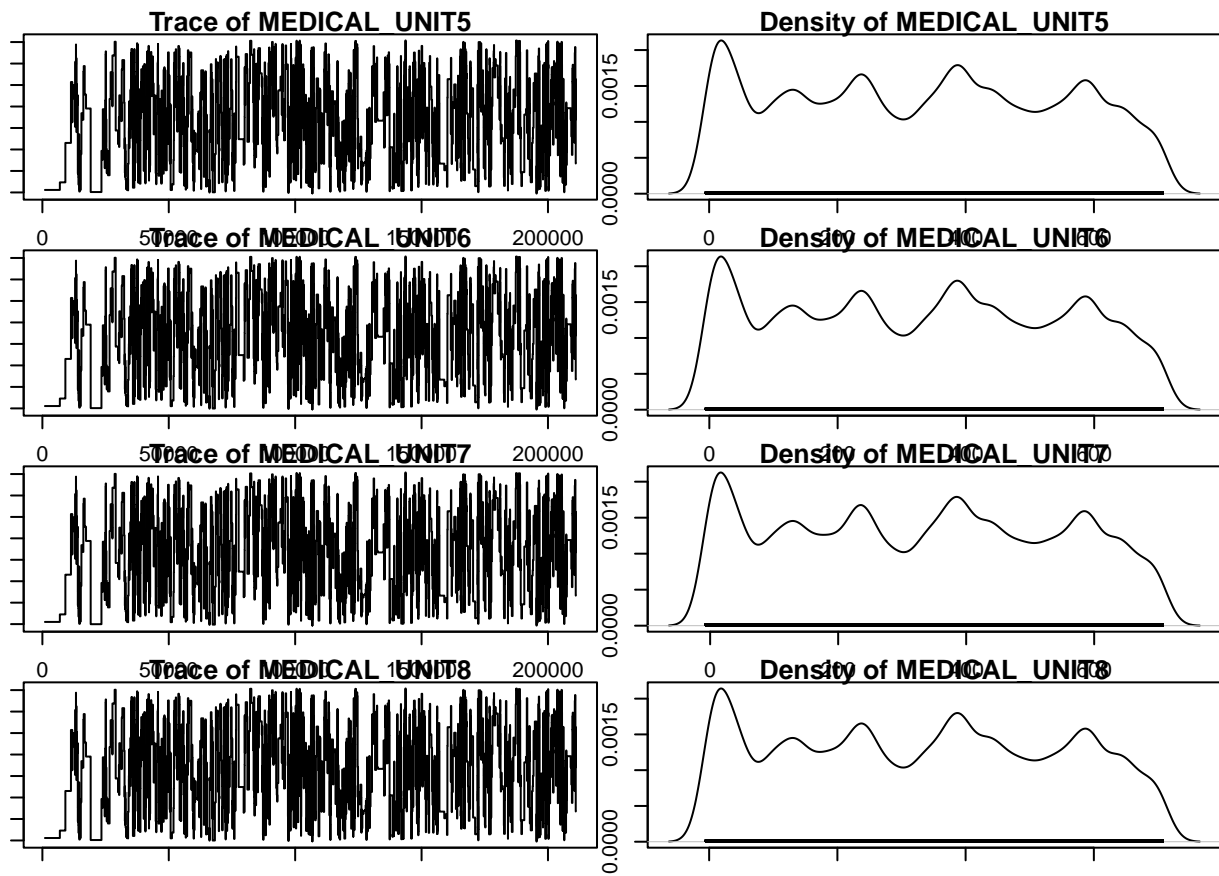
```
library(coda)
library(MASS)
library(MCMCpack)

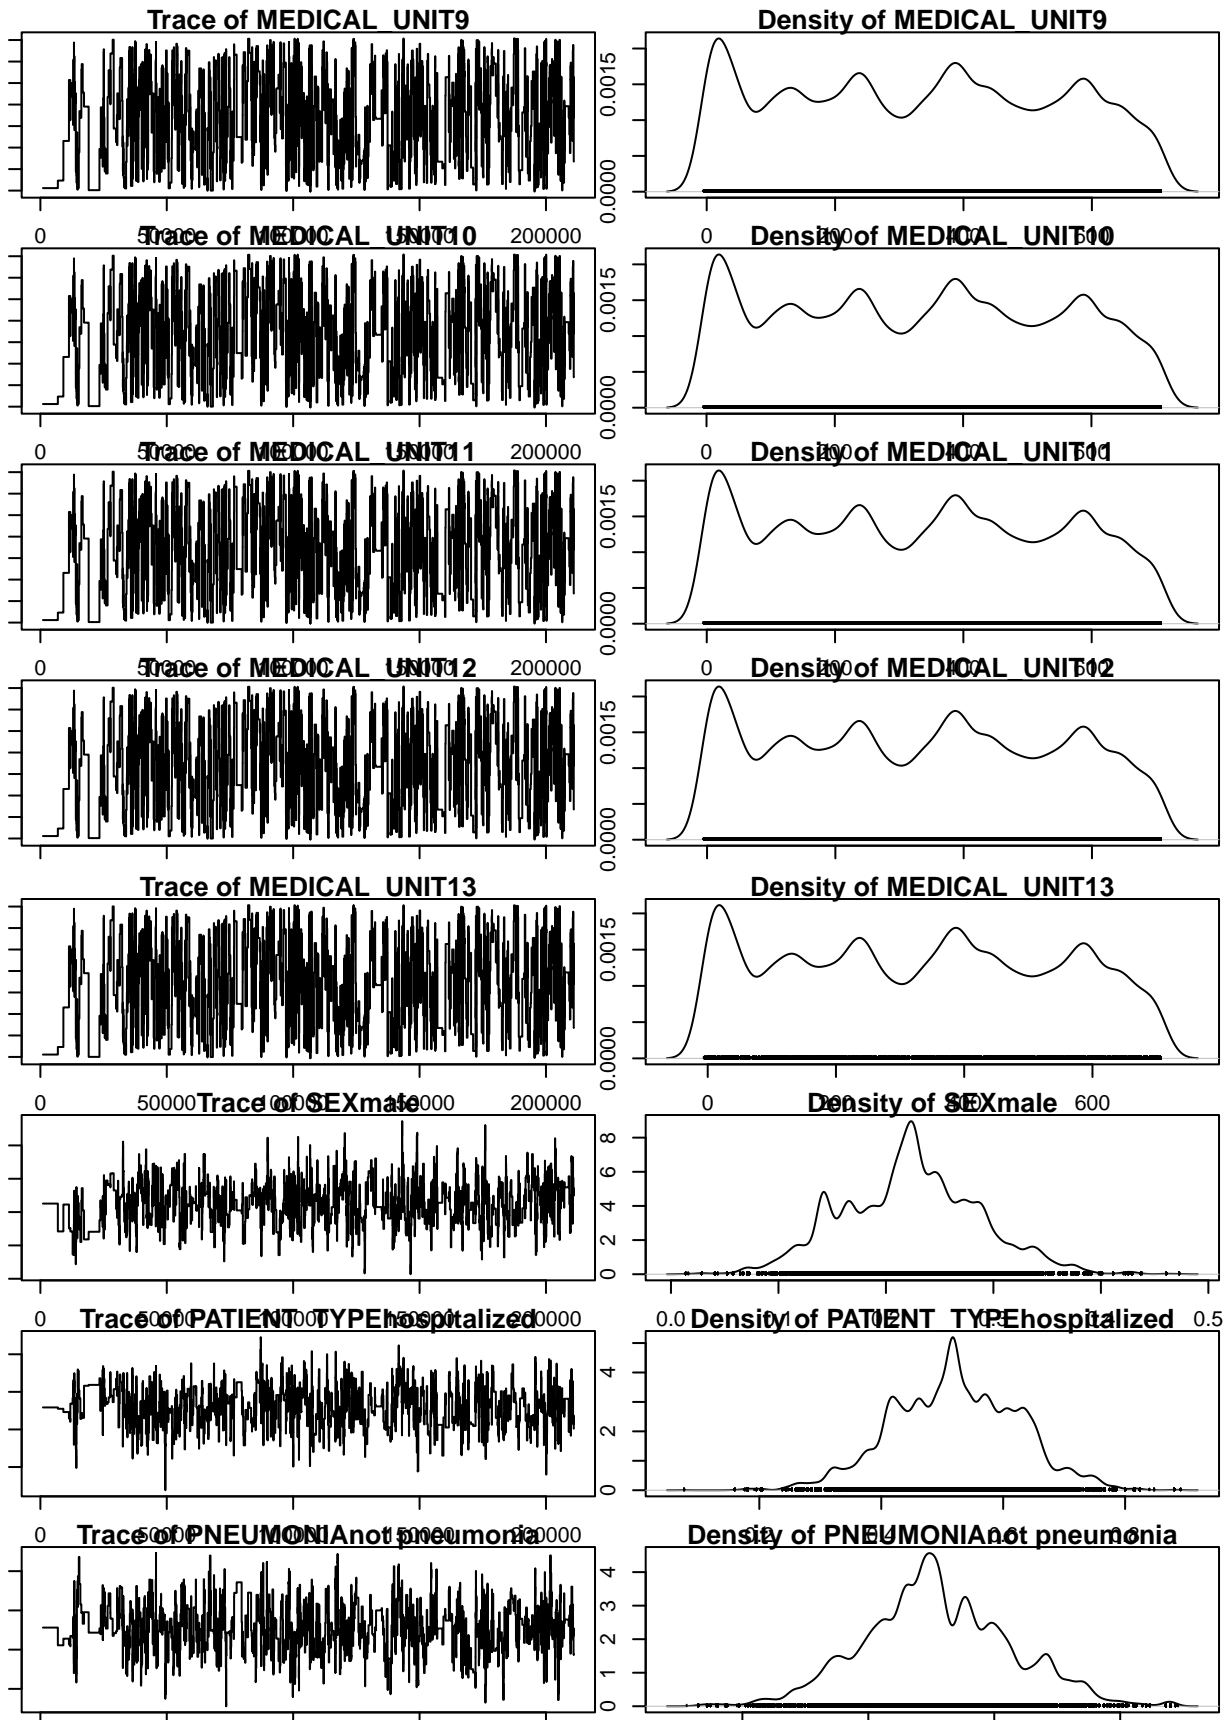
rm(list = setdiff(ls(), c("data", "data.small", "data.test", "data.train")))

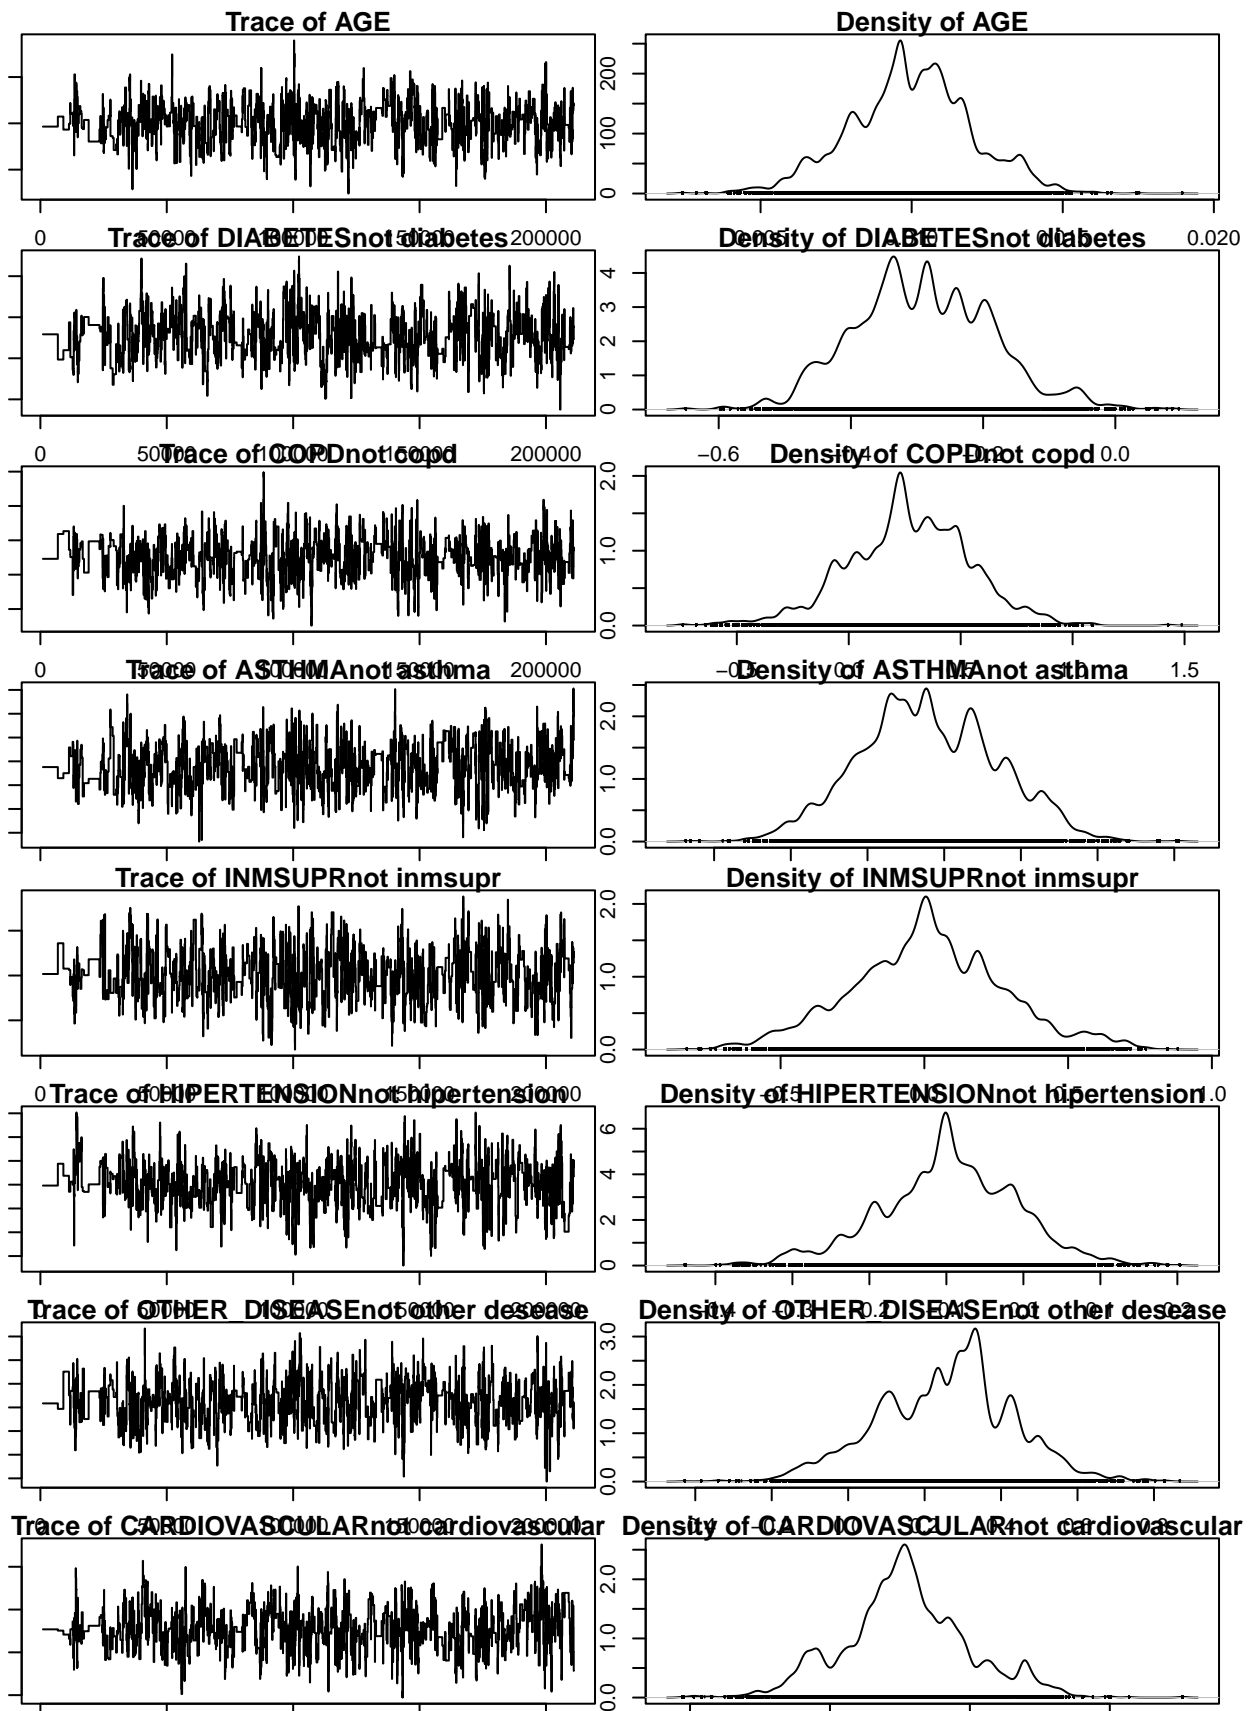
fit = MCMClogit(COVID ~ ., data = data.train, burnin=1000, mcmc=210000)

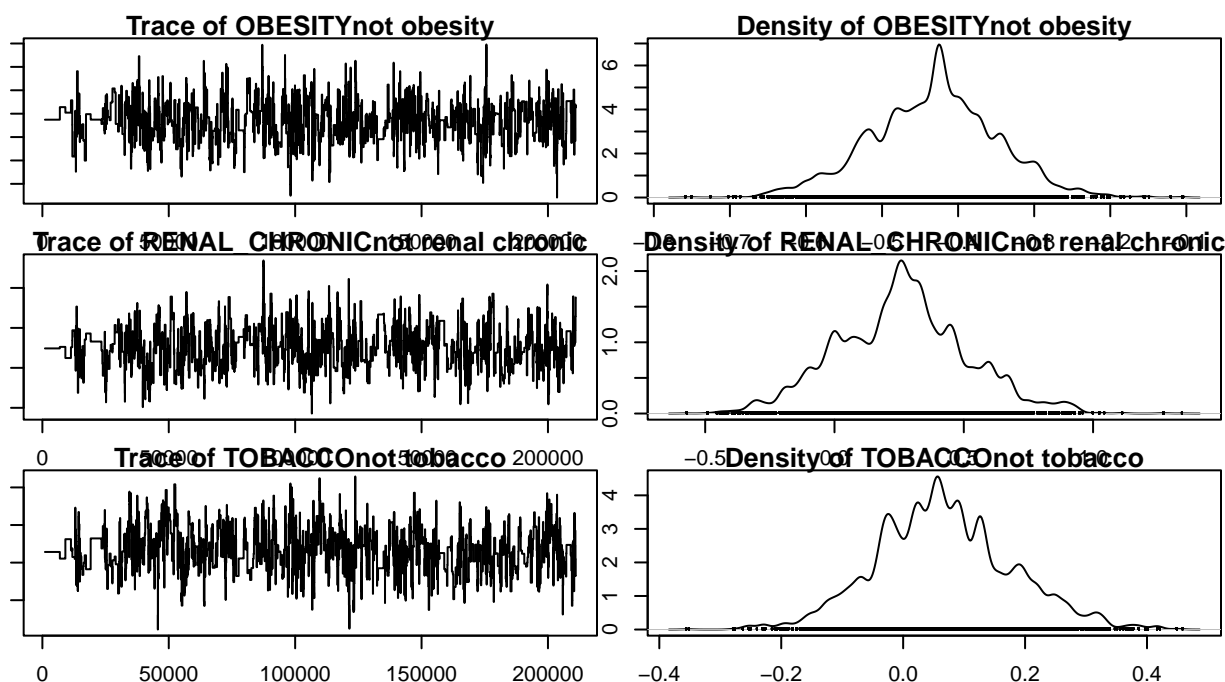
par(mar=c(1, 1, 1, 1))
plot(fit)
```











```
summary(fit)
```

```
##
## Iterations = 1001:211000
## Thinning interval = 1
## Number of chains = 1
## Sample size per chain = 210000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##
```

	Mean	SD	Naive SE	Time-series SE
## (Intercept)	-3.315e+02	2.102e+02	4.586e-01	1.237e+01
## USMERTRUE	-7.082e-03	6.432e-02	1.404e-04	3.512e-03
## MEDICAL_UNIT3	3.305e+02	2.102e+02	4.587e-01	1.237e+01
## MEDICAL_UNIT4	3.302e+02	2.102e+02	4.587e-01	1.237e+01
## MEDICAL_UNIT5	3.304e+02	2.102e+02	4.587e-01	1.237e+01
## MEDICAL_UNIT6	3.297e+02	2.102e+02	4.587e-01	1.237e+01
## MEDICAL_UNIT7	3.285e+02	2.102e+02	4.587e-01	1.237e+01
## MEDICAL_UNIT8	3.302e+02	2.102e+02	4.588e-01	1.238e+01
## MEDICAL_UNIT9	3.304e+02	2.102e+02	4.587e-01	1.237e+01
## MEDICAL_UNIT10	3.310e+02	2.102e+02	4.588e-01	1.237e+01
## MEDICAL_UNIT11	3.305e+02	2.102e+02	4.587e-01	1.237e+01
## MEDICAL_UNIT12	3.301e+02	2.102e+02	4.587e-01	1.237e+01
## MEDICAL_UNIT13	3.297e+02	2.102e+02	4.588e-01	1.237e+01
## SEXmale	2.232e-01	6.345e-02	1.385e-04	3.556e-03
## PATIENT_TYPEhospitalized	5.210e-01	1.025e-01	2.237e-04	5.237e-03
## PNEUMONIAnot pneumonia	-4.900e-01	1.167e-01	2.547e-04	6.279e-03
## AGE	9.984e-03	2.024e-03	4.416e-06	1.055e-04
## DIABETESnot diabetes	-2.905e-01	1.046e-01	2.284e-04	5.764e-03
## COPDnot copd	2.674e-01	2.751e-01	6.003e-04	1.438e-02
## ASTHMAnot asthma	1.584e-01	1.792e-01	3.912e-04	9.317e-03

```

## INMSUPRnot inmsupr      1.221e-02 2.655e-01 5.794e-04      1.455e-02
## HIPERTENSIONnot hypertension -9.758e-02 8.905e-02 1.943e-04      4.602e-03
## OTHER_DISEASEnot other disease 2.542e-01 1.752e-01 3.823e-04      9.469e-03
## CARDIOVASCULARnot cardiovascular 2.721e-01 2.248e-01 4.905e-04      1.171e-02
## OBESITYnot obesity -4.323e-01 8.318e-02 1.815e-04      4.179e-03
## RENAL_CHRONICnot renal chronic 2.617e-01 2.511e-01 5.480e-04      1.377e-02
## TOBACCOnot tobacco      6.633e-02 1.117e-01 2.437e-04      5.882e-03
##
## 2. Quantiles for each variable:
##
##              2.5%      25%      50%      75%
## (Intercept) -6.928e+02 -509.58594 -335.57127 -141.04887
## USMERTRUE -1.320e-01 -0.04972 -0.00610 0.03678
## MEDICAL_UNIT3 2.253e+00 139.79847 333.95952 508.19720
## MEDICAL_UNIT4 1.856e+00 139.42718 333.82917 508.12804
## MEDICAL_UNIT5 1.842e+00 139.84156 334.02260 508.53996
## MEDICAL_UNIT6 1.153e+00 139.06455 333.30264 507.57020
## MEDICAL_UNIT7 -4.698e-01 137.34562 334.23637 505.86906
## MEDICAL_UNIT8 2.030e+00 139.18505 333.50224 508.73075
## MEDICAL_UNIT9 2.033e+00 139.51839 333.93265 508.35535
## MEDICAL_UNIT10 2.444e+00 140.31565 334.39394 508.75437
## MEDICAL_UNIT11 2.335e+00 139.80875 334.21112 508.13608
## MEDICAL_UNIT12 1.771e+00 139.38199 333.72562 508.08403
## MEDICAL_UNIT13 2.485e-01 139.38551 334.27532 509.23532
## SEXmale 1.051e-01 0.17978 0.22426 0.26490
## PATIENT_TYPEhospitalized 3.179e-01 0.45051 0.51730 0.59470
## PNEUMONIAnot pneumonia -7.166e-01 -0.56573 -0.49585 -0.41229
## AGE 5.949e-03 0.00874 0.01001 0.01127
## DIABETESnot diabetes -4.776e-01 -0.36044 -0.29044 -0.21722
## COPDnot copd -2.929e-01 0.09774 0.26009 0.44847
## ASTHMAnot asthma -1.814e-01 0.03977 0.15185 0.28185
## INMSUPRnot inmsupr -5.218e-01 -0.15912 0.01011 0.18092
## HIPERTENSIONnot hypertension -2.976e-01 -0.15084 -0.09844 -0.03817
## OTHER_DISEASEnot other disease -1.019e-01 0.12654 0.27452 0.35631
## CARDIOVASCULARnot cardiovascular -1.511e-01 0.14171 0.26879 0.40693
## OBESITYnot obesity -6.006e-01 -0.48557 -0.42567 -0.37930
## RENAL_CHRONICnot renal chronic -2.090e-01 0.08979 0.26231 0.42746
## TOBACCOnot tobacco -1.453e-01 -0.00921 0.05861 0.12871
##
##              97.5%
## (Intercept) -3.59546
## USMERTRUE 0.12214
## MEDICAL_UNIT3 691.46183
## MEDICAL_UNIT4 690.92586
## MEDICAL_UNIT5 690.62031
## MEDICAL_UNIT6 690.19355
## MEDICAL_UNIT7 688.96992
## MEDICAL_UNIT8 691.17591
## MEDICAL_UNIT9 691.06858
## MEDICAL_UNIT10 691.53135
## MEDICAL_UNIT11 691.04387
## MEDICAL_UNIT12 690.85454
## MEDICAL_UNIT13 689.60722
## SEXmale 0.35069
## PATIENT_TYPEhospitalized 0.72167

```

```
## PNEUMONIAnot pneumonia      -0.25679
## AGE                          0.01387
## DIABETESnot diabetes        -0.06078
## COPDnot copd                0.80971
## ASTHMAnot asthma            0.50358
## INMSUPRnot inmsupr          0.57217
## HIPERTENSIONnot hypertension 0.07425
## OTHER_DISEASEnot other disease 0.57825
## CARDIOVASCULARnot cardiovascular 0.73216
## OBESITYnot obesity          -0.27201
## RENAL_CHRONICnot renal chronic 0.80040
## TOBACCOnot tobacco          0.30335
```

From the Bayesian point of view, we see that the CI for all the parameters does not contain 0 so theoretically all of the predictors are significant with an $\alpha = 5\%$.

Lasso

```
rm(list = setdiff(ls(), c("data", "data.small", "data.test", "data.train")))

library(monomvn)

## Loading required package: pls
##
## Attaching package: 'pls'
## The following object is masked from 'package:caret':
##
##      R2
## The following object is masked from 'package:stats':
##
##      loadings
## Loading required package: lars
## Loaded lars 1.3
##
## Attaching package: 'monomvn'
## The following object is masked from 'package:MCMCpack':
##
##      rwish
x = data.frame(lapply(subset(data.train, select = -c(COVID)), function(x) as.numeric((x))))

adaptTo0And1 = function(col.name, df) {
  index = which(names(df) == col.name)

  if (length(index) != 0) {
    df[, index] = ifelse(df[, index] == 2, 0, df[, index])
  }

  return(df)
}
x = adaptTo0And1("SEX", x)
```

```

x = adaptTo0And1("PATIENT_TYPE", x)
x = adaptTo0And1("PNEUMONIA", x)
x = adaptTo0And1("DIABETES", x)
x = adaptTo0And1("COPD", x)
x = adaptTo0And1("ASTHMA", x)
x = adaptTo0And1("INMSUPR", x)
x = adaptTo0And1("HIPERTENSION", x)
x = adaptTo0And1("OTHER_DISEASE", x)
x = adaptTo0And1("CARDIOVASCULAR", x)
x = adaptTo0And1("OBESITY", x)
x = adaptTo0And1("RENAL_CHRONIC", x)
x = adaptTo0And1("TOBACCO", x)

```

```
summary(x)
```

##	USMER	MEDICAL_UNIT	SEX	PATIENT_TYPE
##	Min. :0.0000	Min. : 2.000	Min. :0.0000	Min. :0.0000
##	1st Qu.:0.0000	1st Qu.: 4.000	1st Qu.:0.0000	1st Qu.:1.0000
##	Median :0.0000	Median :12.000	Median :0.0000	Median :1.0000
##	Mean :0.3535	Mean : 9.017	Mean :0.4893	Mean :0.8286
##	3rd Qu.:1.0000	3rd Qu.:12.000	3rd Qu.:1.0000	3rd Qu.:1.0000
##	Max. :1.0000	Max. :13.000	Max. :1.0000	Max. :1.0000
##	PNEUMONIA	AGE	DIABETES	COPD
##	Min. :0.0000	Min. : 0.00	Min. :0.0000	Min. :0.0000
##	1st Qu.:0.0000	1st Qu.:30.00	1st Qu.:0.0000	1st Qu.:0.0000
##	Median :0.0000	Median :40.00	Median :0.0000	Median :0.0000
##	Mean :0.1258	Mean :41.53	Mean :0.1116	Mean :0.0124
##	3rd Qu.:0.0000	3rd Qu.:52.00	3rd Qu.:0.0000	3rd Qu.:0.0000
##	Max. :1.0000	Max. :98.00	Max. :1.0000	Max. :1.0000
##	ASTHMA	INMSUPR	HIPERTENSION	OTHER_DISEASE
##	Min. :0.00000	Min. :0.0000	Min. :0.0000	Min. :0.00000
##	1st Qu.:0.00000	1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:0.00000
##	Median :0.00000	Median :0.0000	Median :0.0000	Median :0.00000
##	Mean :0.02941	Mean :0.0136	Mean :0.1436	Mean :0.03181
##	3rd Qu.:0.00000	3rd Qu.:0.0000	3rd Qu.:0.0000	3rd Qu.:0.00000
##	Max. :1.00000	Max. :1.0000	Max. :1.0000	Max. :1.00000
##	CARDIOVASCULAR	OBESITY	RENAL_CHRONIC	TOBACCO
##	Min. :0.0000	Min. :0.000	Min. :0.0000	Min. :0.00000
##	1st Qu.:0.0000	1st Qu.:0.000	1st Qu.:0.0000	1st Qu.:0.00000
##	Median :0.0000	Median :0.000	Median :0.0000	Median :0.00000
##	Mean :0.0194	Mean :0.141	Mean :0.0162	Mean :0.08102
##	3rd Qu.:0.0000	3rd Qu.:0.000	3rd Qu.:0.0000	3rd Qu.:0.00000
##	Max. :1.0000	Max. :1.000	Max. :1.0000	Max. :1.00000

```

y = data.train$COVID
fit = blasso(x, y, mprior = c(0,1))

```


```

## t=100, m=8
## t=200, m=11
## t=300, m=9
## t=400, m=9
## t=500, m=9
## t=600, m=11
## t=700, m=9



```

```
## t=800, m=12
## t=900, m=10
plot(fit, burnin=200)
points(drop(fit$b), col=2, pch=20)
points(drop(fit$b), col=3, pch=18)
legend("topleft", c("blasso-map", "lasso", "lsr"),
      col=c(2,2,3), pch=c(21,20,18))
```


Boxplots of regression coefficients



```
plot(fit, burnin=200, which="m")
```



```
acf(fit$m)
```



```
fit = blasso(x, y, mprior = c(0,1), T = 10000, thin = 20)
```



```
## t=100, m=9
## t=200, m=10
## t=300, m=11
## t=400, m=13
## t=500, m=9
## t=600, m=10
## t=700, m=11
## t=800, m=10
## t=900, m=9
## t=1000, m=10
## t=1100, m=9
## t=1200, m=8
## t=1300, m=9
## t=1400, m=10
## t=1500, m=8
```




```
## t=1600, m=9
## t=1700, m=10
## t=1800, m=8
## t=1900, m=8
## t=2000, m=9
## t=2100, m=8
## t=2200, m=11
## t=2300, m=10
## t=2400, m=9
## t=2500, m=10
## t=2600, m=8
## t=2700, m=7
## t=2800, m=6
## t=2900, m=8
## t=3000, m=10
## t=3100, m=8
## t=3200, m=10
## t=3300, m=8
## t=3400, m=8
## t=3500, m=11
## t=3600, m=9
## t=3700, m=8
## t=3800, m=8
## t=3900, m=10
## t=4000, m=10
## t=4100, m=8
## t=4200, m=8
## t=4300, m=11
## t=4400, m=8
## t=4500, m=9
## t=4600, m=9
## t=4700, m=7
## t=4800, m=11
## t=4900, m=9
## t=5000, m=10
## t=5100, m=10
## t=5200, m=12
## t=5300, m=8
## t=5400, m=8
## t=5500, m=9
## t=5600, m=10
## t=5700, m=12
## t=5800, m=7
## t=5900, m=11
## t=6000, m=11
## t=6100, m=11
## t=6200, m=7
## t=6300, m=8
## t=6400, m=9
## t=6500, m=11
## t=6600, m=10
## t=6700, m=11
## t=6800, m=9
## t=6900, m=11
```

```
## t=7000, m=10
## t=7100, m=8
## t=7200, m=11
## t=7300, m=9
## t=7400, m=6
## t=7500, m=12
## t=7600, m=8
## t=7700, m=10
## t=7800, m=10
## t=7900, m=8
## t=8000, m=10
## t=8100, m=10
## t=8200, m=9
## t=8300, m=6
## t=8400, m=10
## t=8500, m=12
## t=8600, m=9
## t=8700, m=8
## t=8800, m=10
## t=8900, m=11
## t=9000, m=10
## t=9100, m=8
## t=9200, m=8
## t=9300, m=9
## t=9400, m=9
## t=9500, m=11
## t=9600, m=9
## t=9700, m=8
## t=9800, m=10
## t=9900, m=12
```

```
plot(fit, burnin=1000, which="m")
```




```
acf(fit$m)
```



```
plot(fit, burnin=1000)
points(drop(fit$b), col=2, pch=20)
points(drop(fit$b), col=3, pch=18)
legend("topleft", c("blasso-map", "lasso", "lsr"),
      col=c(2,2,3), pch=c(21,20,18))
```


Boxplots of regression coefficients



```
s <- summary(fit, burnin=1000)
print(s$bn0) # probability that each beta coef != zero
```



```
##      b.1      b.2      b.3      b.4      b.5      b.6      b.7      b.8
## 0.2758889 0.3170000 0.9957778 1.0000000 0.9995556 0.9998889 0.9200000 0.2965556
##      b.9      b.10     b.11     b.12     b.13     b.14     b.15     b.16
## 0.3027778 0.2557778 0.3836667 0.3832222 0.3726667 1.0000000 0.3285556 0.2840000
```

```
barplot(s$bn0, horiz = TRUE)
```





here we see that the

```
plot(fit, burnin=1000, which="s2")
```



```
plot(fit, burnin=1000, which="lambda2")
```



Final Model

```
rm(list = setdiff(ls(), c("data", "data.small", "data.test", "data.train")))

library(R2OpenBUGS)

model = function() {
}
```