

Lab Bayesian

Andres Navarro

2023-04-25

Dataset Lookup

First of all, the objective of this Case Study will be to show the advantages of Bayesian Statistics for small data sets and the ability to estimated better the posterior parameters.. As it is known, Bayesian Statistics allows as to set up prior believes of our predictors with specific probability distributions. This is really useful when we do not have a lot of data and we have some insights on the data. For this reason I have decided to use a data set with 21 variables and I will be reducing the number of observations to simulate what we are trying to show. This data is about the COVID cases in Mexico and the goal is to predict if a patient has COVID or not.

[[<https://www.kaggle.com/datasets/meirnazri/covid19-dataset>][Dataset]]

```
rm(list = ls())
data = read.csv("data.csv", header = TRUE)
dim(data)
```

```
## [1] 1048575      21
```

```
summary(data)
```

```
##      USMER      MEDICAL_UNIT      SEX      PATIENT_TYPE
## Min.   :1.000   Min.   : 1.000   Min.   :1.000   Min.   :1.000
## 1st Qu.:1.000   1st Qu.: 4.000   1st Qu.:1.000   1st Qu.:1.000
## Median :2.000   Median :12.000   Median :1.000   Median :1.000
## Mean   :1.632   Mean   : 8.981   Mean   :1.499   Mean   :1.191
## 3rd Qu.:2.000   3rd Qu.:12.000   3rd Qu.:2.000   3rd Qu.:1.000
## Max.   :2.000   Max.   :13.000   Max.   :2.000   Max.   :2.000
##  DATE_DIED      INTUBED      PNEUMONIA      AGE
## Length:1048575   Min.   : 1.00   Min.   : 1.000   Min.   : 0.00
## Class :character 1st Qu.:97.00  1st Qu.: 2.000   1st Qu.: 30.00
## Mode  :character Median :97.00  Median : 2.000   Median : 40.00
##                Mean   :79.52  Mean   : 3.347   Mean   : 41.79
##                3rd Qu.:97.00  3rd Qu.: 2.000   3rd Qu.: 53.00
##                Max.   :99.00  Max.   :99.000   Max.   :121.00
##  PREGNANT      DIABETES      COPD      ASTHMA
## Min.   : 1.00   Min.   : 1.000   Min.   : 1.000   Min.   : 1.000
## 1st Qu.: 2.00   1st Qu.: 2.000   1st Qu.: 2.000   1st Qu.: 2.000
## Median :97.00   Median : 2.000   Median : 2.000   Median : 2.000
## Mean   :49.77   Mean   : 2.186   Mean   : 2.261   Mean   : 2.243
## 3rd Qu.:97.00   3rd Qu.: 2.000   3rd Qu.: 2.000   3rd Qu.: 2.000
## Max.   :98.00   Max.   :98.000   Max.   :98.000   Max.   :98.000
##  INMSUPR      HYPERTENSION      OTHER_DISEASE      CARDIOVASCULAR
## Min.   : 1.000   Min.   : 1.000   Min.   : 1.000   Min.   : 1.000
## 1st Qu.: 2.000   1st Qu.: 2.000   1st Qu.: 2.000   1st Qu.: 2.000
```

```
## Median : 2.000 Median : 2.000 Median : 2.000 Median : 2.000
## Mean : 2.298 Mean : 2.129 Mean : 2.435 Mean : 2.262
## 3rd Qu.: 2.000 3rd Qu.: 2.000 3rd Qu.: 2.000 3rd Qu.: 2.000
## Max. :98.000 Max. :98.000 Max. :98.000 Max. :98.000
## OBESITY RENAL_CHRONIC TOBACCO CLASIFFICATION_FINAL
## Min. : 1.000 Min. : 1.000 Min. : 1.000 Min. :1.000
## 1st Qu.: 2.000 1st Qu.: 2.000 1st Qu.: 2.000 1st Qu.:3.000
## Median : 2.000 Median : 2.000 Median : 2.000 Median :6.000
## Mean : 2.125 Mean : 2.257 Mean : 2.214 Mean :5.306
## 3rd Qu.: 2.000 3rd Qu.: 2.000 3rd Qu.: 2.000 3rd Qu.:7.000
## Max. :98.000 Max. :98.000 Max. :98.000 Max. :7.000
## ICU
## Min. : 1.00
## 1st Qu.:97.00
## Median :97.00
## Mean :79.55
## 3rd Qu.:97.00
## Max. :99.00
```

The raw data set consists of 21 unique features and 1,048,576 unique patients. In the Boolean features, 1 means “yes” and 2 means “no”. values as 97 and 99 are missing data.

- sex: 1 for female and 2 for male.
- age: of the patient.
- classification: covid test findings. Values 1-3 mean that the patient was diagnosed with covid in different
- degrees. 4 or higher means that the patient is not a carrier of covid or that the test is inconclusive.
- patient type: type of care the patient received in the unit. 1 for returned home and 2 for hospitalization.
- pneumonia: whether the patient already have air sacs inflammation or not.
- pregnancy: whether the patient is pregnant or not.
- diabetes: whether the patient has diabetes or not.
- copd: Indicates whether the patient has Chronic obstructive pulmonary disease or not.
- asthma: whether the patient has asthma or not.
- inmsupr: whether the patient is immunosuppressed or not.
- hypertension: whether the patient has hypertension or not.
- cardiovascular: whether the patient has heart or blood vessels related disease.
- renal chronic: whether the patient has chronic renal disease or not.
- other disease: whether the patient has other disease or not.
- obesity: whether the patient is obese or not.
- tobacco: whether the patient is a tobacco user.
- usmr: Indicates whether the patient treated medical units of the first, second or third level.
- medical unit: type of institution of the National Health System that provided the care.
- intubed: whether the patient was connected to the ventilator.
- icu: Indicates whether the patient had been admitted to an Intensive Care Unit.
- date died: If the patient died indicate the date of death, and 9999-99-99 otherwise.

Here we can see a summary of the data, first we have to clean and adapt the data so we can work on it. First of all, I will create the variable that we want to predict that is if a patient has been diagnosed with COVID or not.

```
data$COVID = ifelse(data$CLASIFFICATION_FINAL <= 3, 1, 2)
data = subset(data, select = -c(CLASIFFICATION_FINAL))
```

```
convertToLogic = function(col.name, df) {
  index = which(names(df) == col.name)
  print(index)
```

```

if (length(index) != 0) {
  df[, index] = ifelse(df[, index] == 2, 0, df[, index])
  df[, index] = as.logical(df[, index])
}

return(df)
}

```

This column will tell us if a patient has been diagnosed with COVID or not. Then, I will factor and format all the other variables to adapt them properly.

```
data = convertToLogic("COVID", data)
```

```
## [1] 21
```

```

data$USMER = ifelse(data$USMER == 2, 0, data$USMER)
data$USMER = as.logical(data$USMER)

data$MEDICAL_UNIT = factor(data$MEDICAL_UNIT)

data$SEX = factor(data$SEX, labels = c("female", "male"), levels = c(1, 2))

data$PATIENT_TYPE = factor(data$PATIENT_TYPE, labels = c("returned home", "hospitalized"), levels = c(1, 2))

data$INTUBED = factor(data$INTUBED, labels = c("intubed", "not intubed"), levels = c(1, 2))

data$PNEUMONIA = factor(data$PNEUMONIA, labels = c("pneumonia", "not pneumonia"), levels = c(1, 2))

data$PREGNANT = factor(data$PREGNANT, labels = c("pregnant", "not pregnant"), levels = c(1, 2))

data$DIABETES = factor(data$DIABETES, labels = c("diabetes", "not diabetes"), levels = c(1, 2))

data$COPD = factor(data$COPD, labels = c("copd", "not copd"), levels = c(1, 2))

data$ASTHMA = factor(data$ASTHMA, labels = c("asthma", "not asthma"), levels = c(1, 2))

data$INMSUPR = factor(data$INMSUPR, labels = c("inmsupr", "not inmsupr"), levels = c(1, 2))

data$HIPERTENSION = factor(data$HIPERTENSION, labels = c("hipertension", "not hipertension"), levels = c(1, 2))

data$OTHER_DISEASE = factor(data$OTHER_DISEASE, labels = c("other disease", "not other disease"), levels = c(1, 2))

data$CARDIOVASCULAR = factor(data$CARDIOVASCULAR, labels = c("cardiovascular", "not cardiovascular"), levels = c(1, 2))

data$OBESITY = factor(data$OBESITY, labels = c("obesity", "not obesity"), levels = c(1, 2))

data$RENAL_CHRONIC = factor(data$RENAL_CHRONIC, labels = c("renal chronic", "not renal chronic"), levels = c(1, 2))

data$TOBACCO = factor(data$TOBACCO, labels = c("tobacco", "not tobacco"), levels = c(1, 2))

data$ICU = factor(data$ICU, labels = c("icu", "not icu"), levels = c(1, 2))

data = subset(data, select = -c(DATE_DIED))

```

```
summary(data)
```

```
##      USMER      MEDICAL_UNIT      SEX      PATIENT_TYPE
## Mode :logical  12      :602995  female:525064  returned home:848544
## FALSE:662903  4      :314405  male :523511  hospitalized :200031
## TRUE :385672  6      : 40584
##      9      : 38116
##      3      : 19175
##      8      : 10399
##      (Other): 22901
##      INTUBED      PNEUMONIA      AGE
## intubed      : 33656  pneumonia :140038  Min.   : 0.00
## not intubed:159050  not pneumonia:892534  1st Qu.: 30.00
## NA's         :855869  NA's      : 16003  Median : 40.00
##                                     Mean    : 41.79
##                                     3rd Qu.: 53.00
##                                     Max.    :121.00
##
##      PREGNANT      DIABETES      COPD
## pregnant      : 8131  diabetes      :124989  copd      : 15062
## not pregnant:513179  not diabetes:920248  not copd:1030510
## NA's          :527265  NA's          : 3338  NA's      : 3003
##
##
##
##      ASTHMA      INMSUPR      HIPERTENSION
## asthma      : 31572  inmsupr      : 14170  hypertension :162729
## not asthma:1014024  not inmsupr:1031001  not hypertension:882742
## NA's        : 2979  NA's         : 3404  NA's        : 3104
##
##
##
##      OTHER_DISEASE      CARDIOVASCULAR      OBESITY
## other disease : 28040  cardiovascular : 20769  obesity :159816
## not other disease:1015490  not cardiovascular:1024730  not obesity:885727
## NA's          : 5045  NA's          : 3076  NA's      : 3032
##
##
##
##      RENAL_CHRONIC      TOBACCO      ICU
## renal chronic : 18904  tobacco      : 84376  icu      : 16858
## not renal chronic:1026665  not tobacco:960979  not icu:175685
## NA's          : 3006  NA's         : 3220  NA's     :856032
##
##
##
##      COVID
## Mode :logical
## FALSE:656596
## TRUE :391979
```

```
##  
##  
##  
##
```

Here we see that the data is correctly formatted but there are some missing value, so let's fix that.

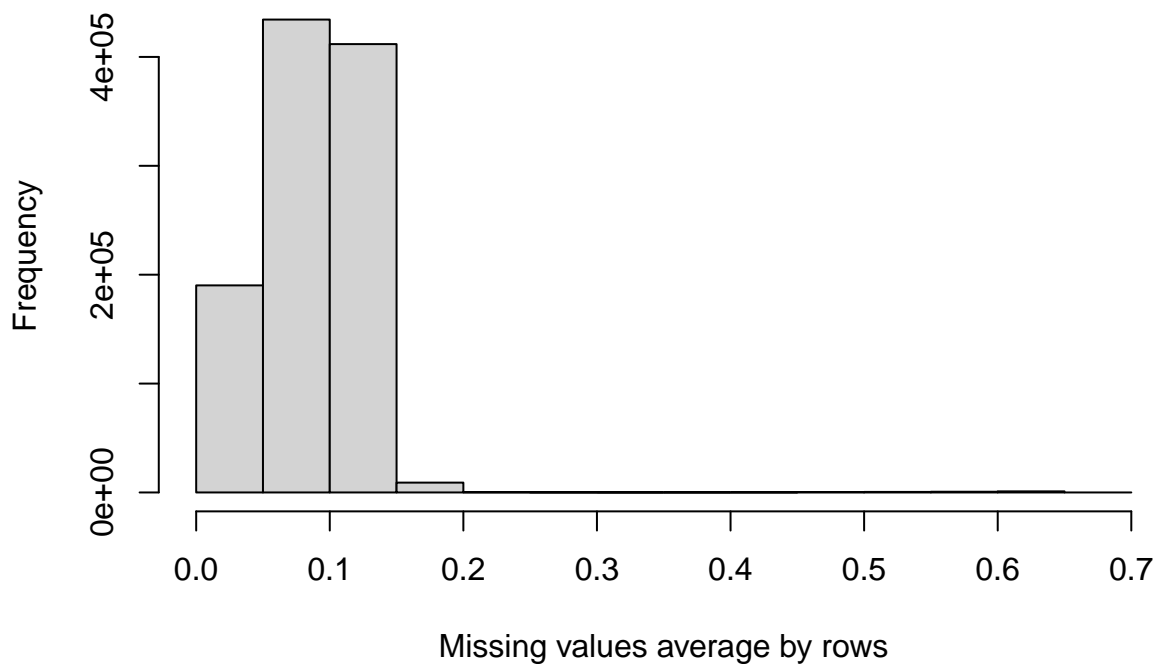
Data Cleaning

First we will see how many missing values there are by rows so we can remove some columns that have a lot of missing values.

```
print(length(which(is.na(data))))
```

```
## [1] 2288376
```

```
hist(rowMeans(is.na(data)), xlab = c("Missing values average by rows"), main = c())
```




Here we see that there are 3 columns with the most missing values.

```
indexesEmptyCols = which(colMeans(is.na(data)) != 0)
```

```
colsWithNA = sort(colMeans(is.na(data[, indexesEmptyCols])),  
                  decreasing = TRUE)
```

```
barplot(colsWithNA, las=2)
```



And the columns that have the most missing values are ICU, INTUBED, and PREGNANT, so let's remove them.

```
data = subset(data, select = -c(ICU, INTUBED, PREGNANT))
```

```
print(length(which(is.na(data))))
```

```
## [1] 49210
```

```
data = na.omit(data)
```

```
length(unique(which(is.na(data))))
```

```
## [1] 0
```

```
summary(data)
```

```
##      USMER      MEDICAL_UNIT      SEX      PATIENT_TYPE
## Mode :logical  12      :591811  female:513216  returned home:833253
## FALSE:658255   4      :307177  male  :511936  hospitalized :191899
## TRUE :366897   6      : 37868
##                      9      : 37384
##                      3      : 18660
##                      8      : 10097
##                      (Other): 22155
##      PNEUMONIA      AGE      DIABETES
## pneumonia      :137599  Min.   : 0.00  diabetes      :122415
## not pneumonia:887553  1st Qu.: 30.00  not diabetes:902737
##                      Median : 40.00
##                      Mean   : 41.89
##                      3rd Qu.: 53.00
##                      Max.   :121.00
##
```

```
##          COPD          ASTHMA          INMSUPR
##  copd      : 14376   asthma      : 30497   inmsupr      : 13588
##  not copd:1010776   not asthma:994655   not inmsupr:1011564
##
##
##
##
##          HIPERTENSION          OTHER_DISEASE
##  hypertension      :159577   other disease      : 27131
##  not hypertension:865575   not other disease:998021
##
##
##
##
##          CARDIOVASCULAR          OBESITY          RENAL_CHRONIC
##  cardiovascular      : 20126   obesity      :156961   renal chronic      : 18351
##  not cardiovascular:1005026   not obesity:868191   not renal chronic:1006801
##
##
##
##
##          TOBACCO          COVID
##  tobacco      : 82675   Mode :logical
##  not tobacco:942477   FALSE:636274
##                      TRUE :388878
##
##
##
##
```

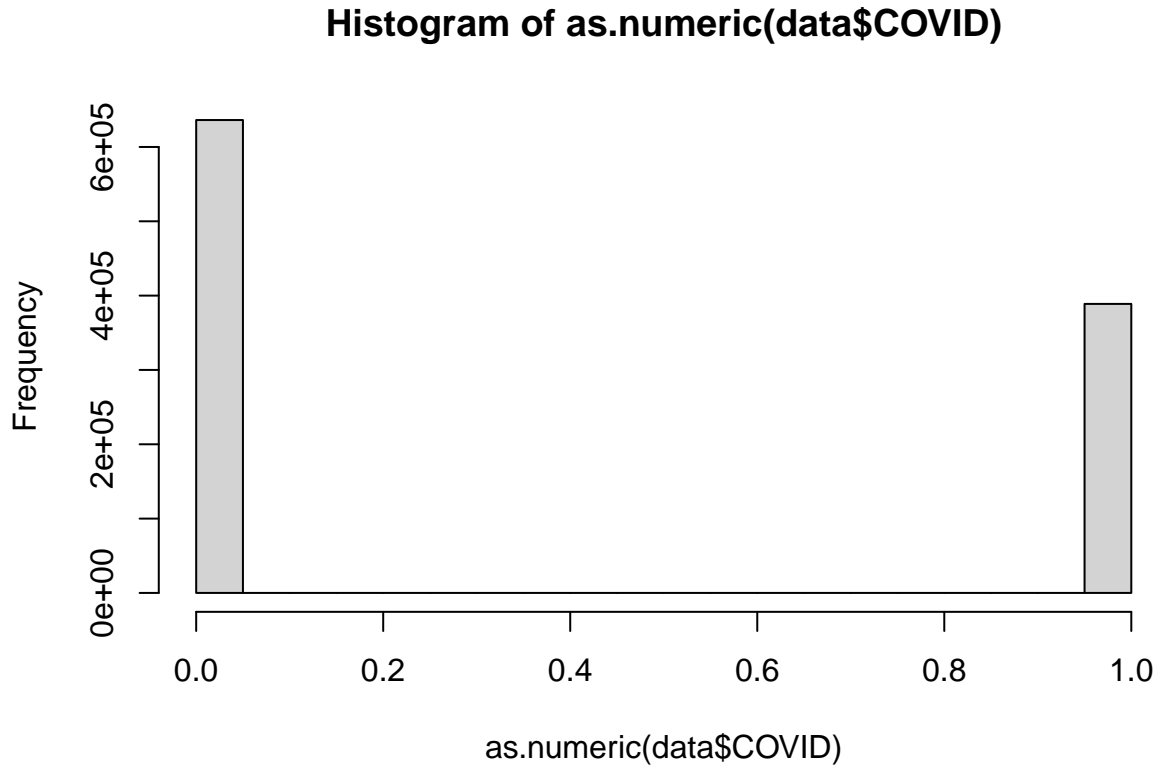
Now the data set is clean so let's start working on it. But first we will shrink it to 1000 observations to work with.

```
data.small = data[sample(nrow(data), size=1000),]
```

Bayesian Analysis of the covid variable

First of all, let's plot a histogram of the COVID variable (the one we want to predict) and see.

```
rm(list = setdiff(ls(), c("data", "data.small")))
hist(as.numeric(data$COVID))
```



This is as we expected as we are going to be predicting a binary variable. So let's use a Bernoulli distribution to explain this data and see how well it fits. First of all let's compute the analytical posterior distribution of the covid variable.

Analytical Study

1. We assume a Bernoulli distribution for COVID, we will use X to denote that variable.

$$X \mid \theta \sim \text{Bernoulli}(\theta)$$

$$f(x \mid \theta) = \theta^x \cdot (1 - \theta)^{1-x}$$

2. As we do not have any prior knowledge on the probability of a patient of having covid, we will define the prior distribution as an improper prior. Moreover, we will be using a Beta distribution as in the end we will get a posterior conjugate which will be much easier to work with.

$$\theta \sim \text{Beta}(0, 0)$$

$$f(\theta \mid 0, 0) = \frac{\theta^{0-1} \cdot (1 - \theta)^{0-1}}{B(0, 0)}$$

3. Now we get the likelihood

$$f(\text{data} \mid \theta) \propto \theta^k \cdot (1 - \theta)^{n-k}$$

Being n the total number of observations and k the positive ones.

4. And finally the posterior distribution

$$f(\theta | data) = \frac{\theta^{k-1} \cdot (1-\theta)^{n-k-1}}{B(k, n-k)}$$

$$\theta | data \sim Beta(k, n-k)$$

So now that we have the posterior distribution let's obtain the prediction of the next value called Y given the data

$$Y | \theta \sim Bernuilli(\theta)$$

$$P(Y = 1|data) = \int_{-\infty}^{\infty} P(Y = 1|\theta) \cdot P(\theta|data)d\theta = \frac{B(k+1, n-k)}{B(k, n-k)}$$

```
n = as.numeric(length(data.small$COVID))
k = as.numeric(length(which(data.small$COVID)))
print(beta(k+1,n-k)/beta(k, n-k))
```

```
## [1] 0.385
```

And here we can see that the probability of a new patient of having covid is 0.362 that is really close to the ML estimator of 0.38

And finally let's try to obtain the same result numerically

Numerical Study

As we know the distribution of the new observation we will get a random sample and compare.

$$Y | \theta \sim Bernuilli(Beta(k, n-k))$$

```
y.sample = rbinom(n, 1, rbeta(1, k, n - k))
mean(y.sample)
```

```
## [1] 0.435
```

Here we see that the estimated probability is almost the same as previously.

```
covid.prob = rbeta(n, k, n - k)
quantile(covid.prob, probs = c(0.025, 0.975))
```

```
##      2.5%      97.5%
## 0.3551356 0.4145273
```

And also we see that the confidence interval for the probability of having covid is pretty narrow, so we can be sure that it is correct.

Data Exploration

Now, we will see if the other variables are useful to predict if a patient has covid or not.


```
rm(list = setdiff(ls(), c("data", "data.small")))

library(ggplot2) # GGally
library(GGally)
```

```
## Registered S3 method overwritten by 'GGally':  
##   method from  
##   +.gg   ggplot2
```

Let's see how the columns distribute with respect to the covid variable.


```
ggplot(data, aes(x=COVID, y=AGE)) +  
  geom_boxplot()
```



Here we see that there is a visible difference between the mean of the covid, so this can be a useful variable to use in our model.


```
plot(table(data$COVID, data$USMER))
```

table(data\$COVID, data\$USMER)




```
plot(table(data$COVID, data$MEDICAL_UNIT))
```

table(data\$COVID, data\$MEDICAL_UNIT)




```
plot(table(data$COVID, data$SEX))
```

table(data\$COVID, data\$SEX)




```
plot(table(data$COVID, data$PATIENT_TYPE))
```

table(data\$COVID, data\$PATIENT_TYPE)




```
plot(table(data$COVID, data$PNEUMONIA))
```

table(data\$COVID, data\$PNEUMONIA)




```
plot(table(data$COVID, data$DIABETES))
```

table(data\$COVID, data\$DIABETES)




```
plot(table(data$COVID, data$COPD))
```

table(data\$COVID, data\$COPD)




```
plot(table(data$COVID, data$ASTHMA))
```

table(data\$COVID, data\$ASTHMA)




```
plot(table(data$COVID, data$INMSUPR))
```

table(data\$COVID, data\$INMSUPR)




```
plot(table(data$COVID, data$HIPERTENSION))
```

table(data\$COVID, data\$HIPERTENSION)




```
plot(table(data$COVID, data$OTHER_DISEASE))
```

table(data\$COVID, data\$OTHER_DISEASE)




```
plot(table(data$COVID, data$CARDIOVASCULAR))
```

table(data\$COVID, data\$CARDIOVASCULAR)




```
plot(table(data$COVID, data$OBESITY))
```


table(data\$COVID, data\$OBESITY)



```
plot(table(data$COVID, data$RENAL_CHRONIC))
```

table(data\$COVID, data\$RENAL_CHRONIC)



```
plot(table(data$COVID, data$TOBACCO))
```

table(data\$COVID, data\$TOBACCO)

	FALSE	TRUE
tobacco		
not tobacco		

Now, the columns that show off the most are: MEDICAL_UNIT, SEX, PATIENT_TYPE, and PNEUMONIA. This makes sense and we will see after if we are confident that there is a visible difference.

Frequentist LM

Now let's implement a simple LM model to see how well we can predict a patient to have covid.

```
rm(list = setdiff(ls(), c("data")))
```

```
library(caret)
```

```
## Loading required package: lattice
```

```
library(lattice)
```

```
data.small = data[sample(nrow(data), size=10000),]
```

```
index.test = createDataPartition(data.small$COVID, p = 0.5, list = FALSE)
```

```
data.test = data.small[index.test,]
```

```
data.train = data.small[-index.test,]
```

```
rm(index.test)
```

Now first of all let's try to use all the variables to try to predict if a patient has covid or not.

```
fit = train(as.factor(COVID) ~ ., data = data.train, method = "glm", family = "binomial")
```

```
summary(fit)
```

```
##
```

```
## Call:
```

```
## NULL
```

```
##
```

```
## Coefficients: (2 not defined because of singularities)
```

```
## Estimate Std. Error z value Pr(>|z|)
```

```
## (Intercept) -0.480154 1.138193 -0.422 0.67313
## USMERTRUE -0.005030 0.065707 -0.077 0.93898
## MEDICAL_UNIT2 NA NA NA NA
## MEDICAL_UNIT3 -0.211763 1.044873 -0.203 0.83939
## MEDICAL_UNIT4 -0.356177 1.026262 -0.347 0.72854
## MEDICAL_UNIT5 -0.383722 1.079604 -0.355 0.72227
## MEDICAL_UNIT6 -0.597059 1.037737 -0.575 0.56506
## MEDICAL_UNIT7 -12.453065 228.461441 -0.055 0.95653
## MEDICAL_UNIT8 -0.012526 1.075717 -0.012 0.99071
## MEDICAL_UNIT9 -0.685330 1.039360 -0.659 0.50965
## MEDICAL_UNIT10 -0.100638 1.074661 -0.094 0.92539
## MEDICAL_UNIT11 0.695247 1.103862 0.630 0.52880
## MEDICAL_UNIT12 -0.468333 1.025802 -0.457 0.64799
## MEDICAL_UNIT13 NA NA NA NA
## SEXmale 0.189070 0.060842 3.108 0.00189 **
## PATIENT_TYPEhospitalized 0.415423 0.100113 4.150 3.33e-05 ***
## `PNEUMONIA`not pneumonia` -0.628042 0.110615 -5.678 1.36e-08 ***
## AGE 0.009659 0.002056 4.699 2.62e-06 ***
## `DIABETES`not diabetes` -0.061639 0.100342 -0.614 0.53902
## `COPD`not copd` 0.429007 0.280563 1.529 0.12624
## `ASTHMA`not asthma` -0.062139 0.177057 -0.351 0.72562
## `INMSUPR`not inmsupr` 0.573871 0.253999 2.259 0.02386 *
## `HIPERTENSION`not hypertension` -0.123096 0.093500 -1.317 0.18799
## `OTHER_DISEASE`not other disease` -0.142194 0.170069 -0.836 0.40310
## `CARDIOVASCULAR`not cardiovascular` -0.002049 0.213432 -0.010 0.99234
## `OBESITY`not obesity` -0.267177 0.083974 -3.182 0.00146 **
## `RENAL_CHRONIC`not renal chronic` -0.014886 0.210786 -0.071 0.94370
## `TOBACCO`not tobacco` -0.024984 0.112226 -0.223 0.82383
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 6644.6 on 4998 degrees of freedom
## Residual deviance: 6346.9 on 4973 degrees of freedom
## AIC: 6398.9
##
## Number of Fisher Scoring iterations: 11
```

Here we see that there are a lot of variables that are useless. As the p value of the betas is really high for most of them.

```
confusionMatrix(as.factor(data.test$COVID), predict(fit, newdata = data.test))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction FALSE TRUE
##      FALSE  2820  275
##      TRUE   1422  484
##
##           Accuracy : 0.6607
##           95% CI : (0.6474, 0.6738)
##      No Information Rate : 0.8482
##      P-Value [Acc > NIR] : 1
```

```
##
##           Kappa : 0.1867
##
## Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.6648
##           Specificity : 0.6377
##           Pos Pred Value : 0.9111
##           Neg Pred Value : 0.2539
##           Prevalence : 0.8482
##           Detection Rate : 0.5639
##           Detection Prevalence : 0.6189
##           Balanced Accuracy : 0.6512
##
##           'Positive' Class : FALSE
##
```

Here, we see that we get an accuracy of 0.6578 so it is not that bad, probably it is because we only have a few significant variables as we saw in the correlation graph. So let's try a simpler model.

```
fit = train(as.factor(COVID) ~ USMER + PNEUMONIA + MEDICAL_UNIT + DIABETES + HIPERTENSION + AGE + PATIENT_TYPE)
summary(fit)
```

```
##
## Call:
## NULL
##
## Coefficients: (2 not defined because of singularities)
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      0.212780   1.039192   0.205    0.838
## USMERTRUE        -0.002279   0.065414  -0.035    0.972
## `PNEUMONIAnot pneumonia` -0.640955   0.109895  -5.832 5.46e-09 ***
## MEDICAL_UNIT2              NA          NA      NA      NA
## MEDICAL_UNIT3      -0.240153   1.038209  -0.231    0.817
## MEDICAL_UNIT4      -0.371058   1.019501  -0.364    0.716
## MEDICAL_UNIT5      -0.435579   1.073237  -0.406    0.685
## MEDICAL_UNIT6      -0.615138   1.031161  -0.597    0.551
## MEDICAL_UNIT7     -12.346076  228.299334  -0.054    0.957
## MEDICAL_UNIT8        0.058285   1.068378   0.055    0.956
## MEDICAL_UNIT9      -0.714162   1.032705  -0.692    0.489
## MEDICAL_UNIT10     -0.128067   1.068324  -0.120    0.905
## MEDICAL_UNIT11       0.620423   1.097352   0.565    0.572
## MEDICAL_UNIT12     -0.486438   1.019184  -0.477    0.633
## MEDICAL_UNIT13              NA          NA      NA      NA
## `DIABETESnot diabetes`   -0.072495   0.099186  -0.731    0.465
## `HIPERTENSIONnot hypertension` -0.140312   0.091580  -1.532    0.125
## AGE                  0.009169   0.002033   4.511 6.46e-06 ***
## PATIENT_TYPEhospitalized  0.417164   0.098949   4.216 2.49e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 6644.6  on 4998  degrees of freedom
```

```
## Residual deviance: 6375.2 on 4982 degrees of freedom
## AIC: 6409.2
##
## Number of Fisher Scoring iterations: 11
```

Now it is better but the medical unit for example, it is only relevant the level 2 and also for other.

```
confusionMatrix(as.factor(data.test$COVID), predict(fit, newdata = data.test))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction FALSE TRUE
##      FALSE  2809  286
##      TRUE   1432  474
##
##           Accuracy : 0.6565
##           95% CI : (0.6431, 0.6696)
##      No Information Rate : 0.848
##      P-Value [Acc > NIR] : 1
##
##           Kappa : 0.1767
##
##  Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.6623
##           Specificity : 0.6237
##           Pos Pred Value : 0.9076
##           Neg Pred Value : 0.2487
##           Prevalence : 0.8480
##           Detection Rate : 0.5617
##      Detection Prevalence : 0.6189
##           Balanced Accuracy : 0.6430
##
##           'Positive' Class : FALSE
##
```

Here we see that the accuracy is almost the same and the kappa so we have not lost a lot of info.

Bayesian LM

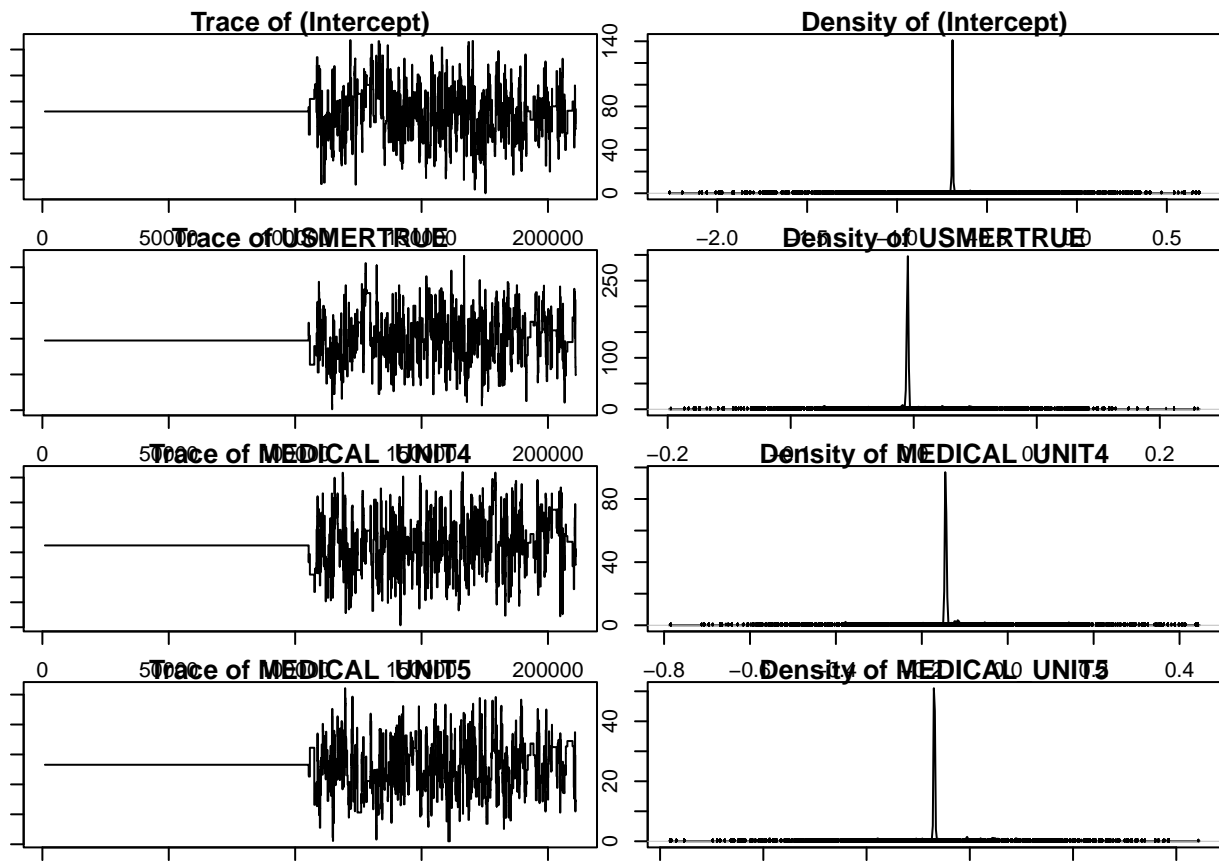
The frequentest approach is easier but we if we want to compute confidence intervals for the parameters or predictive intervals we cannot do them. That is why we will be using the Bayesian approach to better study the effects of each variable with covid and get more conclusions. The power of the Bayesian approach is that we obtain the posterior distribution of the parameters so we can study better the relation and the significance. So let's start.

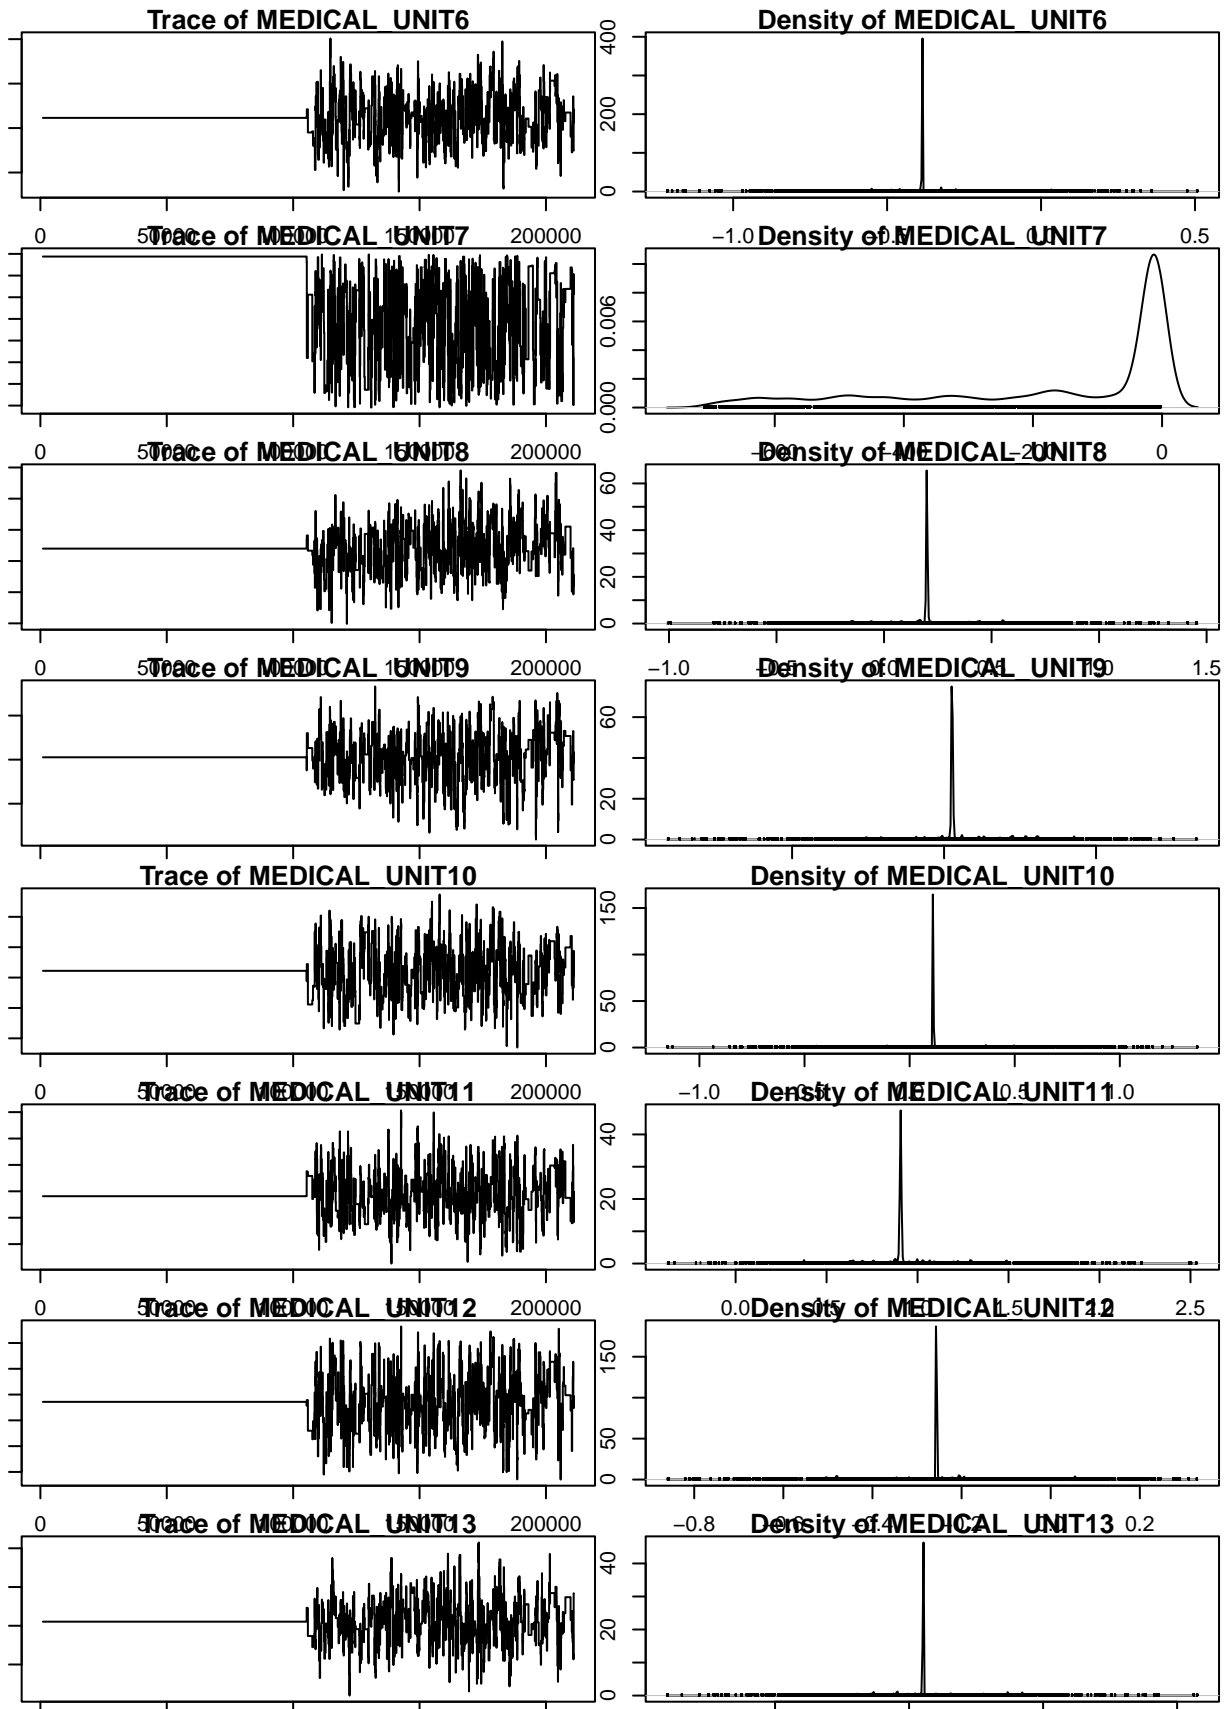
```
library(coda)
library(MASS)
library(MCMCpack)
```


```
rm(list = setdiff(ls(), c("data", "data.small", "data.test", "data.train")))

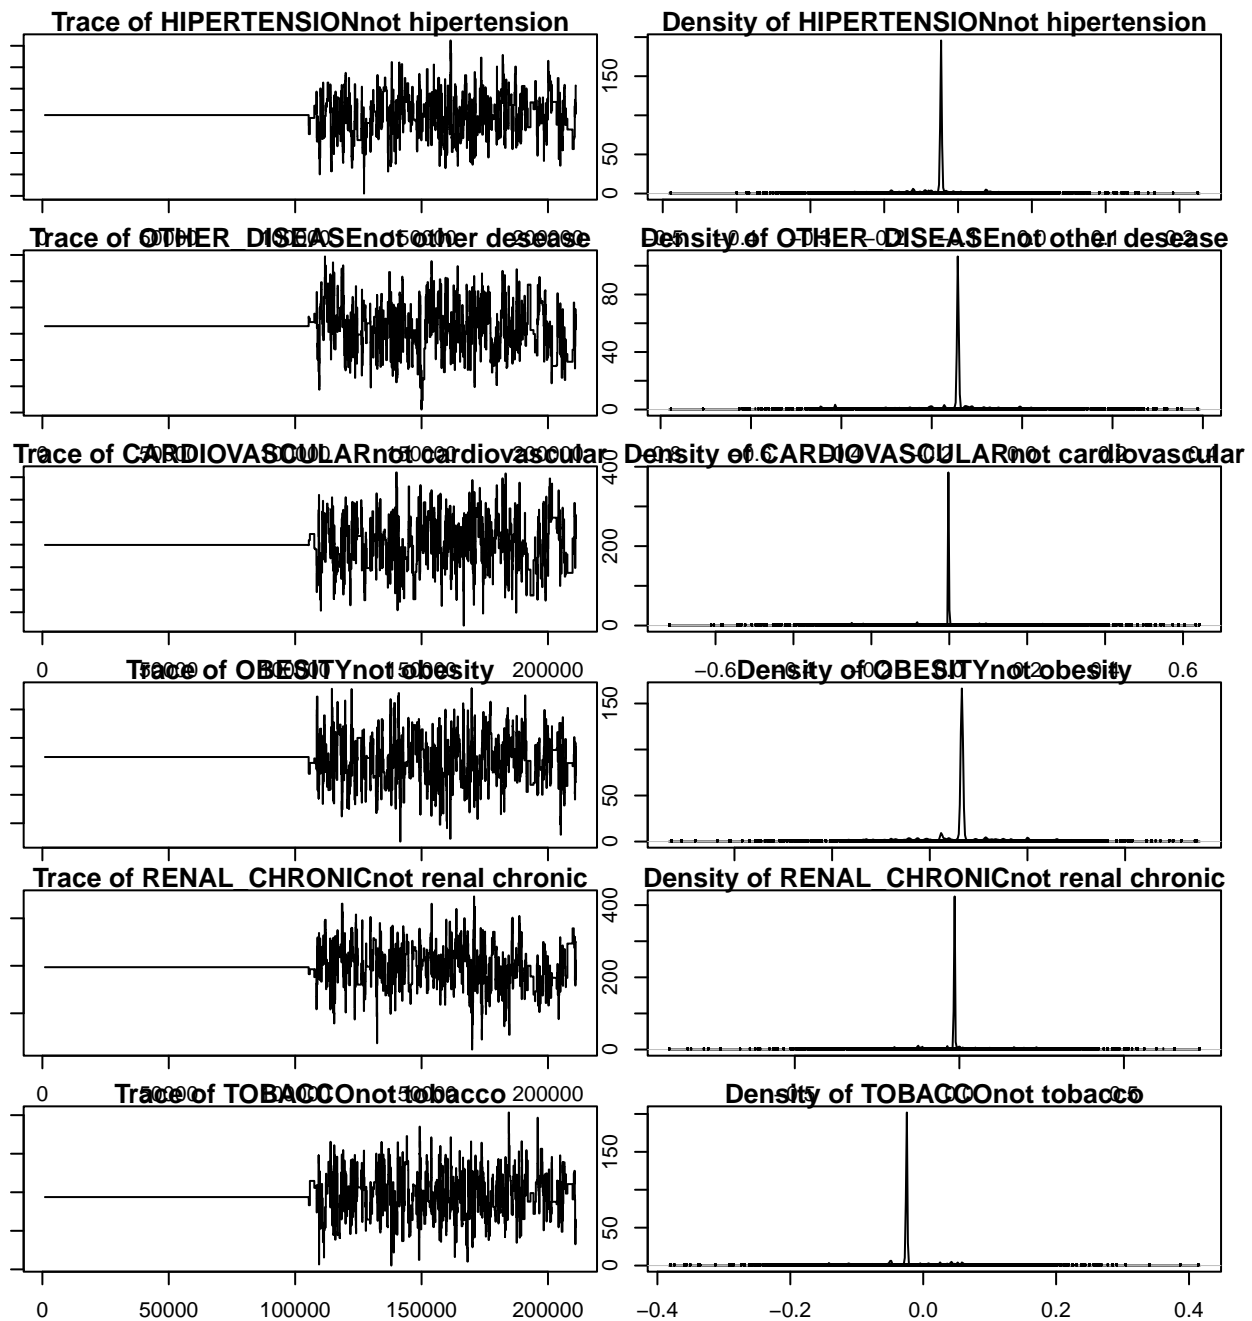
fit = MCMClogit(COVID ~ ., data = data.train, burnin=1000, mcmc=210000)
```

```
par(mar=c(1, 1, 1, 1))
plot(fit)
```









```
summary(fit)
```

```
##
## Iterations = 1001:211000
## Thinning interval = 1
## Number of chains = 1
## Sample size per chain = 210000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##
```

	Mean	SD	Naive SE	Time-series SE
--	------	----	----------	----------------

## (Intercept)	-6.998e-01	3.507e-01	7.652e-04	1.582e-02
## USMERTRUE	-2.014e-03	4.888e-02	1.067e-04	2.415e-03
## MEDICAL_UNIT4	-1.401e-01	1.520e-01	3.316e-04	7.147e-03
## MEDICAL_UNIT5	-1.908e-01	3.086e-01	6.734e-04	1.594e-02
## MEDICAL_UNIT6	-3.831e-01	1.842e-01	4.019e-04	9.073e-03
## MEDICAL_UNIT7	-1.718e+02	2.116e+02	4.618e-01	1.614e+01
## MEDICAL_UNIT8	1.917e-01	2.671e-01	5.829e-04	1.258e-02
## MEDICAL_UNIT9	-4.694e-01	1.975e-01	4.309e-04	9.614e-03
## MEDICAL_UNIT10	1.046e-01	3.017e-01	6.584e-04	1.614e-02
## MEDICAL_UNIT11	9.325e-01	3.048e-01	6.650e-04	1.418e-02
## MEDICAL_UNIT12	-2.571e-01	1.532e-01	3.344e-04	7.306e-03
## MEDICAL_UNIT13	2.185e-01	7.839e-01	1.711e-03	3.979e-02
## SEXmale	1.867e-01	4.212e-02	9.192e-05	1.872e-03
## PATIENT_TYPEhospitalized	4.202e-01	7.253e-02	1.583e-04	3.485e-03
## PNEUMONIAnot pneumonia	-6.254e-01	7.648e-02	1.669e-04	3.583e-03
## AGE	9.627e-03	1.452e-03	3.168e-06	6.652e-05
## DIABETESnot diabetes	-5.706e-02	7.986e-02	1.743e-04	4.401e-03
## COPDnot copd	4.137e-01	2.140e-01	4.671e-04	1.092e-02
## ASTHMAnot asthma	-5.204e-02	1.399e-01	3.052e-04	7.199e-03
## INMSUPRnot inmsupr	5.896e-01	1.794e-01	3.915e-04	8.445e-03
## HIPERTENSIONnot hypertension	-1.247e-01	6.771e-02	1.477e-04	3.301e-03
## OTHER_DISEASEnot other disease	-1.537e-01	1.334e-01	2.911e-04	6.837e-03
## CARDIOVASCULARnot cardiovascular	-5.646e-04	1.553e-01	3.388e-04	7.378e-03
## OBESITYnot obesity	-2.727e-01	5.751e-02	1.255e-04	2.696e-03
## RENAL_CHRONICnot renal chronic	-1.491e-02	1.508e-01	3.291e-04	7.544e-03
## TOBACCOnot tobacco	-1.963e-02	7.927e-02	1.730e-04	3.538e-03

##

2. Quantiles for each variable:

##

##	2.5%	25%	50%	75%
## (Intercept)	-1.513e+00	-7.097e-01	-0.691917	-0.691917
## USMERTRUE	-1.140e-01	-5.030e-03	-0.005030	0.002826
## MEDICAL_UNIT4	-4.895e-01	-1.444e-01	-0.144414	-0.123138
## MEDICAL_UNIT5	-9.499e-01	-2.056e-01	-0.171959	-0.171959
## MEDICAL_UNIT6	-8.155e-01	-3.922e-01	-0.385296	-0.385296
## MEDICAL_UNIT7	-6.527e+02	-3.224e+02	-12.241302	-12.241302
## MEDICAL_UNIT8	-4.164e-01	1.662e-01	0.199237	0.199237
## MEDICAL_UNIT9	-9.487e-01	-4.736e-01	-0.473567	-0.451139
## MEDICAL_UNIT10	-5.816e-01	1.111e-01	0.111125	0.127023
## MEDICAL_UNIT11	2.551e-01	9.070e-01	0.907010	0.950678
## MEDICAL_UNIT12	-6.073e-01	-2.566e-01	-0.256570	-0.246485
## MEDICAL_UNIT13	-1.641e+00	1.638e-01	0.211763	0.211763
## SEXmale	8.943e-02	1.842e-01	0.189070	0.189070
## PATIENT_TYPEhospitalized	2.495e-01	4.154e-01	0.415423	0.421572
## PNEUMONIAnot pneumonia	-7.964e-01	-6.280e-01	-0.628042	-0.620672
## AGE	6.225e-03	9.492e-03	0.009659	0.009659
## DIABETESnot diabetes	-2.261e-01	-6.164e-02	-0.061639	-0.039939
## COPDnot copd	-1.056e-01	3.998e-01	0.429007	0.429007
## ASTHMAnot asthma	-3.494e-01	-6.214e-02	-0.062139	-0.034258
## INMSUPRnot inmsupr	1.962e-01	5.739e-01	0.573871	0.608353
## HIPERTENSIONnot hypertension	-2.826e-01	-1.358e-01	-0.123096	-0.123096
## OTHER_DISEASEnot other disease	-4.616e-01	-1.644e-01	-0.142194	-0.142194
## CARDIOVASCULARnot cardiovascular	-3.526e-01	-2.196e-03	-0.002049	0.004678
## OBESITYnot obesity	-4.138e-01	-2.808e-01	-0.267177	-0.267177

```
## RENAL_CHRONICnot renal chronic -3.498e-01 -1.489e-02 -0.014886 -0.009515
## TOBACCOnot tobacco -1.971e-01 -2.498e-02 -0.024984 -0.013233
## 97.5%
## (Intercept) 0.11840
## USMERTRUE 0.12475
## MEDICAL_UNIT4 0.22249
## MEDICAL_UNIT5 0.45819
## MEDICAL_UNIT6 0.05562
## MEDICAL_UNIT7 -12.24130
## MEDICAL_UNIT8 0.82512
## MEDICAL_UNIT9 -0.05104
## MEDICAL_UNIT10 0.77907
## MEDICAL_UNIT11 1.68145
## MEDICAL_UNIT12 0.11564
## MEDICAL_UNIT13 1.98000
## SEXmale 0.27739
## PATIENT_TYPEhospitalized 0.58735
## PNEUMONIAnot pneumonia -0.42738
## AGE 0.01281
## DIABETESnot diabetes 0.12814
## COPDnot copd 0.90239
## ASTHMAnot asthma 0.27950
## INMSUPRnot inmsupr 1.01206
## HIPERTENSIONnot hipertension 0.03583
## OTHER_DISEASEnot other disease 0.16216
## CARDIOVASCULARnot cardiovascular 0.36474
## OBESITYnot obesity -0.14612
## RENAL_CHRONICnot renal chronic 0.34865
## TOBACCOnot tobacco 0.16941
```

From the Bayesian point of view, we see that the CI for all the parameters does not contain 0 so theoretically all of the predictors are significant with an $\alpha = 5\%$.

Lasso

```
rm(list = setdiff(ls(), c("data", "data.small", "data.test", "data.train")))

library(monomvn)

x = data.frame(lapply(subset(data.train, select = -c(COVID)), function(x) as.numeric((x))))

adaptTo0And1 = function(col.name, df) {
  index = which(names(df) == col.name)

  if (length(index) != 0) {
    df[, index] = ifelse(df[, index] == 2, 0, df[, index])
  }

  return(df)
}

x = adaptTo0And1("SEX", x)
x = adaptTo0And1("PATIENT_TYPE", x)
x = adaptTo0And1("PNEUMONIA", x)
x = adaptTo0And1("DIABETES", x)
```

```

x = adaptTo0And1("COPD", x)
x = adaptTo0And1("ASTHMA", x)
x = adaptTo0And1("INMSUPR", x)
x = adaptTo0And1("HIPERTENSION", x)
x = adaptTo0And1("OTHER_DISEASE", x)
x = adaptTo0And1("CARDIOVASCULAR", x)
x = adaptTo0And1("OBESITY", x)
x = adaptTo0And1("RENAL_CHRONIC", x)
x = adaptTo0And1("TOBACCO", x)

summary(x)

```

```

##      USMER      MEDICAL_UNIT      SEX      PATIENT_TYPE
## Min.   :0.0000 Min.   : 3.000 Min.   :0.0000 Min.   :0.0000
## 1st Qu.:0.0000 1st Qu.: 4.000 1st Qu.:0.0000 1st Qu.:1.0000
## Median :0.0000 Median :12.000 Median :1.0000 Median :1.0000
## Mean   :0.3493 Mean   : 8.901 Mean   :0.5049 Mean   :0.8088
## 3rd Qu.:1.0000 3rd Qu.:12.000 3rd Qu.:1.0000 3rd Qu.:1.0000
## Max.   :1.0000 Max.   :13.000 Max.   :1.0000 Max.   :1.0000
##      PNEUMONIA      AGE      DIABETES      COPD
## Min.   :0.0000 Min.   : 0.00 Min.   :0.0000 Min.   :0.0000
## 1st Qu.:0.0000 1st Qu.:30.00 1st Qu.:0.0000 1st Qu.:0.0000
## Median :0.0000 Median :41.00 Median :0.0000 Median :0.0000
## Mean   :0.1362 Mean   :42.19 Mean   :0.1236 Mean   :0.0124
## 3rd Qu.:0.0000 3rd Qu.:53.00 3rd Qu.:0.0000 3rd Qu.:0.0000
## Max.   :1.0000 Max.   :99.00 Max.   :1.0000 Max.   :1.0000
##      ASTHMA      INMSUPR      HIPERTENSION      OTHER_DISEASE
## Min.   :0.00000 Min.   :0.000 Min.   :0.0000 Min.   :0.00000
## 1st Qu.:0.00000 1st Qu.:0.000 1st Qu.:0.0000 1st Qu.:0.00000
## Median :0.00000 Median :0.000 Median :0.0000 Median :0.00000
## Mean   :0.03041 Mean   :0.016 Mean   :0.1616 Mean   :0.03361
## 3rd Qu.:0.00000 3rd Qu.:0.000 3rd Qu.:0.0000 3rd Qu.:0.00000
## Max.   :1.00000 Max.   :1.000 Max.   :1.0000 Max.   :1.00000
##      CARDIOVASCULAR      OBESITY      RENAL_CHRONIC      TOBACCO
## Min.   :0.0000 Min.   :0.000 Min.   :0.0000 Min.   :0.00000
## 1st Qu.:0.0000 1st Qu.:0.000 1st Qu.:0.0000 1st Qu.:0.00000
## Median :0.0000 Median :0.000 Median :0.0000 Median :0.00000
## Mean   :0.0212 Mean   :0.153 Mean   :0.0222 Mean   :0.07942
## 3rd Qu.:0.0000 3rd Qu.:0.000 3rd Qu.:0.0000 3rd Qu.:0.00000
## Max.   :1.0000 Max.   :1.000 Max.   :1.0000 Max.   :1.00000

```

```

y = data.train$COVID
fit = blasso(x, y, mprior = c(0,1))

```



```

## t=100, m=7
## t=200, m=7
## t=300, m=13
## t=400, m=9
## t=500, m=7
## t=600, m=8
## t=700, m=10
## t=800, m=9
## t=900, m=7


```

After training the model lets check for stability

```
plot(fit, burnin=200, which="m")
```



```
acf(fit$m)
```



Here we see that there is some autocorrelation and it is not stable enough so to make sure let's add a lot more samples and some thinning.



```
set.seed(111)
fit = blasso(x, y, mprior = c(0,1), T = 10000, thin = 20)
```

```
## t=100, m=9
## t=200, m=10
## t=300, m=7
## t=400, m=11
## t=500, m=8
## t=600, m=10
## t=700, m=9
## t=800, m=6
## t=900, m=9
```


```
## t=1000, m=8
## t=1100, m=6
## t=1200, m=7
## t=1300, m=10
## t=1400, m=10
## t=1500, m=9
## t=1600, m=9
## t=1700, m=8
## t=1800, m=11
## t=1900, m=10
## t=2000, m=8
## t=2100, m=8
## t=2200, m=8
## t=2300, m=9
## t=2400, m=11
## t=2500, m=10
## t=2600, m=10
## t=2700, m=9
## t=2800, m=11
## t=2900, m=10
## t=3000, m=10
## t=3100, m=8
## t=3200, m=10
## t=3300, m=7
## t=3400, m=8
## t=3500, m=7
## t=3600, m=9
## t=3700, m=10
## t=3800, m=8
## t=3900, m=8
## t=4000, m=10
## t=4100, m=8
## t=4200, m=9
## t=4300, m=9
## t=4400, m=8
## t=4500, m=10
## t=4600, m=10
## t=4700, m=9
## t=4800, m=9
## t=4900, m=8
## t=5000, m=8
## t=5100, m=10
## t=5200, m=8
## t=5300, m=10
## t=5400, m=11
## t=5500, m=10
## t=5600, m=8
## t=5700, m=10
## t=5800, m=12
## t=5900, m=8
## t=6000, m=7
## t=6100, m=12
## t=6200, m=11
## t=6300, m=12
```

```
## t=6400, m=9
## t=6500, m=10
## t=6600, m=10
## t=6700, m=9
## t=6800, m=8
## t=6900, m=8
## t=7000, m=8
## t=7100, m=9
## t=7200, m=10
## t=7300, m=8
## t=7400, m=9
## t=7500, m=8
## t=7600, m=7
## t=7700, m=8
## t=7800, m=10
## t=7900, m=10
## t=8000, m=8
## t=8100, m=10
## t=8200, m=7
## t=8300, m=13
## t=8400, m=9
## t=8500, m=8
## t=8600, m=10
## t=8700, m=9
## t=8800, m=7
## t=8900, m=8
## t=9000, m=8
## t=9100, m=8
## t=9200, m=12
## t=9300, m=9
## t=9400, m=10
## t=9500, m=7
## t=9600, m=11
## t=9700, m=9
## t=9800, m=8
## t=9900, m=11
```

```
plot(fit, burnin=1000, which="m")
```





```
acf(fit$m)
```





Now we do not see any periodicity and it is much stable.

```
plot(fit, burnin=1000, which="s2")
```

```
plot(fit, burnin=1000, which="lambda2")
```




And also we see that it is stable so we can trust that it has converged. So let's see the most important variables.

```

plot(fit, burnin=1000)
points(drop(fit$b), col=2, pch=20)
points(drop(fit$b), col=3, pch=18)
legend("topleft", c("blasso-map", "lasso", "lsr"),
      col=c(2,2,3), pch=c(21,20,18))

```

Boxplots of regression coefficients



```

s <- summary(fit, burnin=1000)
print(s$bn0) # probability that each beta coef != zero

```

```


##      b.1      b.2      b.3      b.4      b.5      b.6      b.7      b.8
## 0.2672222 0.5098889 0.9303333 1.0000000 1.0000000 0.9997778 0.3573333 0.4225556
##      b.9      b.10     b.11     b.12     b.13     b.14     b.15     b.16
## 0.2615556 0.6800000 0.5247778 0.3250000 0.2566667 0.9750000 0.2667778 0.2620000

```

```

barplot(s$bn0, horiz = TRUE)

```



Here we see that the most important variables are SEX, PATIENT_TYPE, PNEUMONIA, AGE, CARDIOVASCULAR. All in all we can conclude that the most important variables to predict if a patient has covid or not is:

- Sex: we will see but a specific gender is prone to have covid.
- Patient type: this makes sense because if a patient has been hospitalized it is more likely that it will have covid.
- Age: elder people are more likely to have covid.
- Pneumonia: as the covid makes similar symptoms as pneumonia, patients with it are much more likely to have it.
- Cardiovascular: we can see that covid could be more likely to be contracted if you have cardiovascular problems.

Final Model

Finally we will create a model with the variables selected with lasso.

```
rm(list = setdiff(ls(), c("data", "data.small", "data.test", "data.train")))

library(R2OpenBUGS)

logit.bayes <- function(){
  for( i in 1 : n ) {
    COVID.bin[i] ~ dbern(p[i])
    logit(p[i]) <- b0 + b1 * AGE[i] + b2*SEX.male[i] + b3*PATIENT_TYPE.hospitalized[i] + b4*PNEUMONIA[i]
  }

  b0 ~ dnorm(0.0, 1.0E-6)
  b1 ~ dnorm(0.0, 1.0E-6)
  b2 ~ dnorm(0.0, 1.0E-6)
  b3 ~ dnorm(0.0, 1.0E-6)
  b4 ~ dnorm(0.0, 1.0E-6)
  b5 ~ dnorm(0.0, 1.0E-6)
}

COVID.bin=ifelse(data.train$COVID,1,0)
```

```

n=length(COVID.bin)
SEX.male = ifelse(data.train$SEX=="male",1,0)
PATIENT_TYPE.hospitalized = ifelse(data.train$PATIENT_TYPE=="hospitalized",1,0)
PNEUMONIA.yes = ifelse(data.train$PNEUMONIA=="pneumonia",1,0)
CARDIOVASCULAR.yes = ifelse(data.train$CARDIOVASCULAR=="cardiovascular",1,0)

data <- list(n=n, COVID.bin=COVID.bin, AGE=data.train$AGE, SEX.male = SEX.male, PATIENT_TYPE.hospitalized = PATIENT_TYPE.hospitalized)


inits <- function(){
  list(b0 = 1, b1 = 0, b2 = 0, b3 = 0, b4 = 0, b5 = 0)
}
output <- bugs(data = data, inits = inits, parameters.to.save = c("b0", "b1", "b2", "b3", "b4", "b5"), n.iter = 1000)

output

## Inference for Bugs model at "/tmp/RtmpUmGfKp/model1154db567ac377.txt",
## Current: 1 chains, each with 1000 iterations (first 100 discarded)
## Cumulative: n.sims = 900 iterations saved
##
##      mean sd   2.5%   25%   50%   75%  97.5%
## b0      -1.2 0.1  -1.4  -1.3  -1.2  -1.2  -1.0
## b1       0.0 0.0   0.0   0.0   0.0   0.0   0.0
## b2       0.2 0.1   0.1   0.1   0.2   0.2   0.3
## b3       0.5 0.1   0.3   0.4   0.5   0.5   0.7
## b4       0.6 0.1   0.4   0.6   0.6   0.7   0.9
## b5       0.0 0.2  -0.4  -0.2   0.0   0.1   0.4
## deviance 6396.5 3.6 6392.0 6394.0 6396.0 6398.0 6405.0
##
## DIC info (using the rule, pD = Dbar-Dhat)
## pD = 6.1 and DIC = 6403.0
## DIC is an estimate of expected predictive error (lower deviance is better).


b0.post <-output$sims.list$b0
b1.post <-output$sims.list$b1
b2.post <-output$sims.list$b2
b3.post <-output$sims.list$b3
b4.post <-output$sims.list$b4
b5.post <-output$sims.list$b5
ts.plot(b0.post)

```




```
acf(b0.post)
```

Series b0.post




```
ts.plot(b1.post)
```




```
acf(b1.post)
```

Series b1.post




```
ts.plot(b2.post)
```




```
acf(b2.post)
```

Series b2.post




```
ts.plot(b3.post)
```




```
acf(b3.post)
```

Series b3.post




```
ts.plot(b4.post)
```





```
acf(b4.post)
```

Series b4.post




```
ts.plot(b5.post)
```



```
acf(b5.post)
```

Series b5.post



We see that the model is stable so let's check some assumptions. First let's create the baseline.

```
linear = b0.post
pred.baseline = exp(linear)/(1+exp(linear))

mean(pred.baseline)
```

```
## [1] 0.2303748
```

```
quantile(pred.baseline,c(0.025,0.975))
```

```
##      2.5%      97.5%
```

```
## 0.1997193 0.2643609
```

Here we see that a baseline person has a 0.22 probability of having covid. So now let's compare it to other groups.

Age

20 years old

```
linear = b0.post+b1.post * 20
```

```
pred.prob = exp(linear)/(1+exp(linear))
```

```
mean(pred.prob)
```

```
## [1] 0.2698846
```

```
quantile(pred.prob,c(0.025,0.975))
```

```
##      2.5%      97.5%
```

```
## 0.2469362 0.2938482
```

50 years old

```
linear = b0.post+b1.post * 50
```

```
pred.prob = exp(linear)/(1+exp(linear))
```

```
mean(pred.prob)
```

```
## [1] 0.3368903
```

```
quantile(pred.prob,c(0.025,0.975))
```

```
##      2.5%      97.5%
```

```
## 0.3149719 0.3574110
```

80 years old

```
linear = b0.post+b1.post * 80
```

```
pred.prob = exp(linear)/(1+exp(linear))
```

```
mean(pred.prob)
```

```
## [1] 0.4113386
```

```
quantile(pred.prob,c(0.025,0.975))
```

```
##      2.5%      97.5%
```

```
## 0.3692872 0.4539434
```

Here we can clearly see that the higher the age, the more probability people have to have covid.

Sex

Female

```
linear = b0.post
```

```
pred.prob = exp(linear)/(1+exp(linear))
```

```

mean(pred.prob)

## [1] 0.2303748
quantile(pred.prob,c(0.025,0.975))

##      2.5%      97.5%
## 0.1997193 0.2643609

Male

linear = b0.post+b2.post * 1
pred.prob = exp(linear)/(1+exp(linear))

mean(pred.prob)

## [1] 0.262799
quantile(pred.prob,c(0.025,0.975))

##      2.5%      97.5%
## 0.2311018 0.3004539

```

We see some indication that the male population has higher probability to have covid but it is not significant.

Patient Type

Not hospitalized

```

linear = b0.post
pred.prob = exp(linear)/(1+exp(linear))

mean(pred.prob)

## [1] 0.2303748
quantile(pred.prob,c(0.025,0.975))

##      2.5%      97.5%
## 0.1997193 0.2643609

```

Hospitalized

```

linear = b0.post+b3.post * 1
pred.prob = exp(linear)/(1+exp(linear))

mean(pred.prob)

## [1] 0.322236
quantile(pred.prob,c(0.025,0.975))

##      2.5%      97.5%
## 0.2669285 0.3834891

```

Now, this is really significant and it makes sense. If a patient has been hospitazlied, it is really likely that he/she has covid.

Pneumonia

Not pneumonia

```
linear = b0.post
pred.prob = exp(linear)/(1+exp(linear))

mean(pred.prob)
```

```
## [1] 0.2303748
```

```
quantile(pred.prob,c(0.025,0.975))
```

```
##      2.5%      97.5%
## 0.1997193 0.2643609
```

Pneumonia

```
linear = b0.post+b4.post * 1
pred.prob = exp(linear)/(1+exp(linear))

mean(pred.prob)
```

```
## [1] 0.3608047
```

```
quantile(pred.prob,c(0.025,0.975))
```

```
##      2.5%      97.5%
## 0.3014263 0.4228176
```

Same conclusions for pneumonia

Cardiovascular

Not cardiovascular

```
linear = b0.post
pred.prob = exp(linear)/(1+exp(linear))

mean(pred.prob)
```

```
## [1] 0.2303748
```

```
quantile(pred.prob,c(0.025,0.975))
```

```
##      2.5%      97.5%
## 0.1997193 0.2643609
```

Cardiovascular

```
linear = b0.post+b5.post * 1
pred.prob = exp(linear)/(1+exp(linear))

mean(pred.prob)
```

```
## [1] 0.2265248
```

```
quantile(pred.prob,c(0.025,0.975))
```

```
##      2.5%      97.5%
## 0.1544751 0.3139523
```

But for cardiovascular we are not really sure that it plays a big role, so we will not make assumptions.

Conclusions

We have seen that many variables in this data set are not useful to predict if a patient has covid or not, but with the Bayesian approach and lasso, we have found some that are significant (in this case: age, patient type, and pneumonia) this makes sense. But the power of Bayesian approach is that we obtain the posterior distribution so we can see how significant and if we trust the variable. As for example, the sex and the cardiovascular had really big confidence intervals so we discarded them. This shows the usefulness of the Bayesian approach in comparison with the frequentest one.